

Databases HW4

12.2

1. $\sigma_{a < 50,000}(R)$: Direct access should be better, since the B+ tree will have an additional lookup cost.
2. $\sigma_{a = 50,000}(R)$: Linear hashed index
3. $\sigma_{a > 50,000 \wedge a < 50,000}(R)$: B+ tree
4. $\sigma_{a \neq 50,000}(R)$: Direct access, since will require scanning through entire set most likely

12.6

with R_t as total tuples in R , R_p as pages in R , same for S

1. Index Nested Hash: $R_p + R_t * (1.2 + 1)$
2. Index Clustered B+ Tree: $R_p + R_t * (3 + 1)$
3. Index Unclustered B+ Tree: $R_p + R_t * (3 + R_t / S_t)$
4. Sorted Merge Joins: $5 * (R_t + S_t)$

15.7

1. $\pi_{D.dname, F.budget} ($
 $((\pi_{E.did}(\sigma_{E.sal \geq 59000, E.hobby = 'yodeling'}(E))) \bowtie_{did} (\pi_{D.did, D.dname}(\sigma_{D.floor = 1}(D))))$
 $\bowtie_{did} (\pi_{F.budget, F.did}(F)))$
2. DEF and DFE. Starting with EF is not considered, since it's a cross product, and we only start with D since we're doing left deep joins.
3.
 - a. employees: $50000 * 1 / 50 * 1/200 = 5$ (assuming even distributions); departments: $5000 / 2 = 2500$, finances: 5000
 - b. ((DE)F)

16.5

1. Very basically, decreasing the use of **X** locks by replacing them with other locks that can be shared will always increase concurrency. For example, if transaction one consists of **Inc A, Inc B** and transaction two consists of **Inc B, Inc A**, typical databases management would require **X** locks on each action and potentially results in deadlock, whereas this system would proceed smoothly with no issues.
2. In the worst case scenario, **I** and **D** locks are just the same as **X** locks, and since strict 2PL guarantees serializability for even the most exclusive type locks, it should intuitively work for **I** and **D** as well.

17.4

1. Strict 2PL with deadlock prevention
 - S1: T1 gets S lock on X; T2 tries to get an X lock on X, fails and aborts; T3 gets X lock on Y; T1 tries X lock on Y, waits; T3 finishes, commits, releases; T1 gets X lock, finishes, commits, releases; T2 finishes.
 - S2: Same scenario as S1, except T2 acquires a X lock Y initially then aborts
2. Strict 2PL with deadlock detection
 - S1: T1 gets S lock on X; T2 waits for X lock on X; T3 gets X lock on Y; T1 waits for X lock on Y; T3 finishes, commits, releases; T1 gets X lock on Y, finishes and releases; T2 gets X lock on X, Y, finishes
 - S2: T1 gets an S lock on X, T2 gets an X lock on Y, T3 waits for Y, T1 waits for Y, T2 waits for X. Deadlock.
3. Conservative
 - S1: T1 gets S on X, X on Y, commits, releases, same with T2, same with T3.
 - S2: Same as S1
4. Optimistic
 - S1: T1 validates and commits fine, T2 aborted and restarted, same with T3
 - S2: same as S1
5. Timestamp conc control + Thomas Write Rule
 - both sequences complete in that order
6. Multiversion concurrency control
 - T1 reads X, T2 writes X, T2 writes Y, T3 writes Y, T1 fails to write Y, aborted and restarted
 - same as S1