

Introduction to Database Systems

Jennie Duggan
Winter 2016

Support

- Instructor office hours: Monday 2-4pm in Ford 3.221
- TA hours
 - Peer mentors: Diane Liu, Nevil George, Thurs – 5-7pm in T-Lab
 - Grad TA: Dipendra Jha, MW - 5-7pm in Wilkinson Lab
- Discussion group
 - Nevil George, Friday – 1-2pm, location TBA
- Additional peer mentors: Katherine Lin, Shannon Nachreiner
- Use Piazza

Administrivia

- Grades:
 - 20% problem sets
 - 30% labs
 - 25% midterm
 - 25% final
- See Canvas for more details
- Textbook: Database Management Systems (R. Ramakrishnan, J. Gehrke)

Policies

- Full syllabus on Canvas
- Problem sets & programming assignments due at midnight on specified date
- Individual responsibility to keep up-to-date
 - Pro tip: Canvas is the authoritative info source
- Late days
- Don't cheat!

Why databases?

- Used to be the province of banks and retailers
- Seeing a renaissance from:
 - Web
 - Diversity of data already collected/available
 - Cheap storage!

Class Focus: Relational database management systems (RDBMSs)

Key Database Issues

- Database design and querying
- Data analysis
- Concurrency control
- Scalability and Efficiency

Background

- Ted Codd invented the relational model in the seminal paper “A Relational Model of Data for Large Shared Data Banks”
 - Main concept: *relation* = a table with rows and columns.
 - Every relation has a *schema*, which describes the columns.
- Prior 1970, no standard data model.
 - Network model used by Codasyl
 - Hierarchical model used by IMS
- After 1970, IBM built System R as proof-of-concept for relational model and used **SQL** as the query language. SQL eventually became a standard.

Example of a Traditional Database Application


Suppose we are building a system to store the information about:

- students
- courses
- professors
- who takes what, who teaches what

Can we do it without a DBMS ?

Sure we can! Start by storing the data in files:

students.txt courses.txt professors.txt

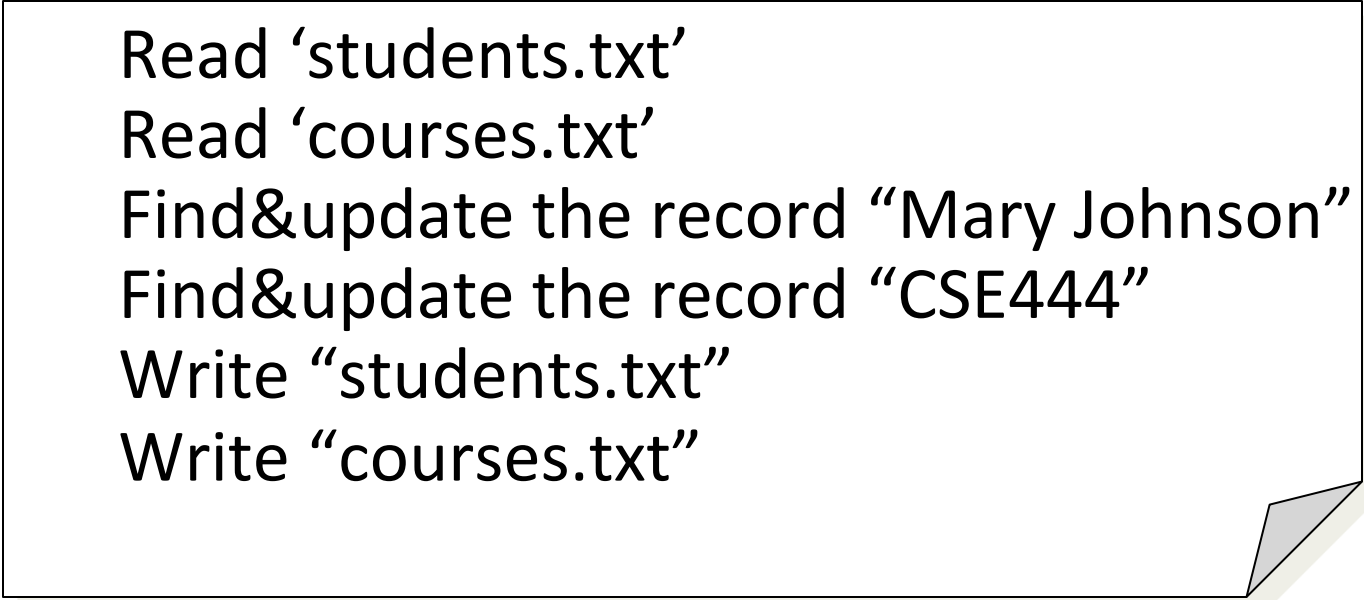


Now write C or Java programs to implement specific tasks

Doing it without a DBMS...

- Enroll “Mary Johnson” in “CSE444”:

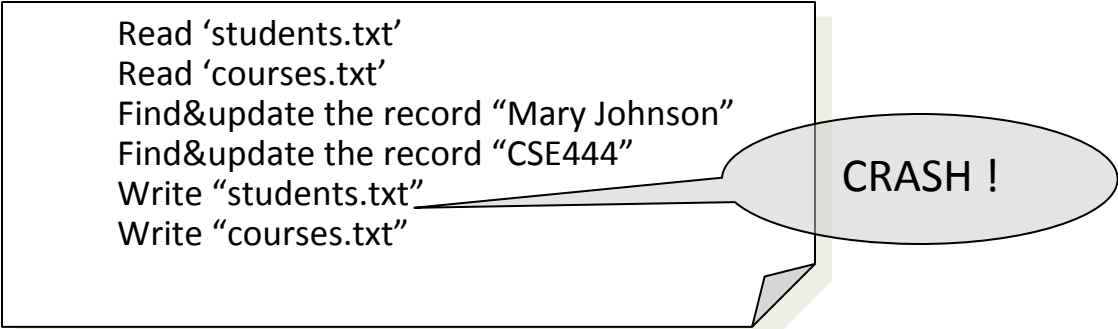
Write a C program to do the following:



- Read ‘students.txt’
- Read ‘courses.txt’
- Find&update the record “Mary Johnson”
- Find&update the record “CSE444”
- Write “students.txt”
- Write “courses.txt”

Problems without a DBMS...

- System crashes:



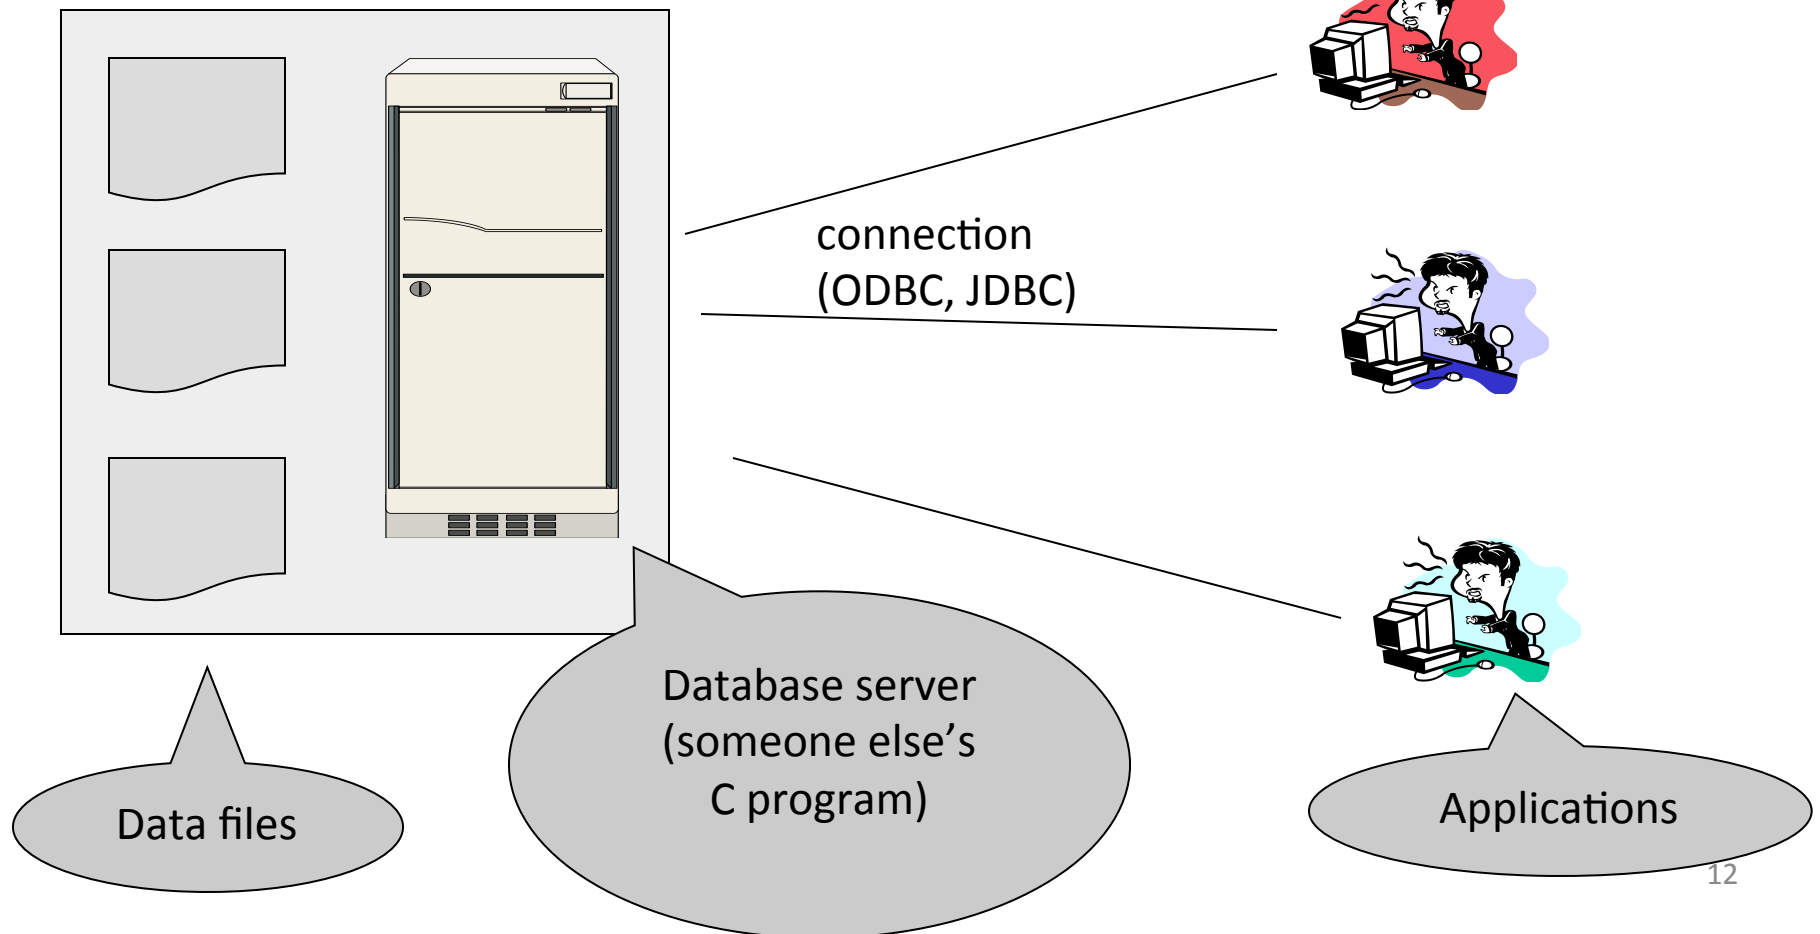
```
Read 'students.txt'  
Read 'courses.txt'  
Find&update the record "Mary Johnson"  
Find&update the record "CSE444"  
Write "students.txt"  
Write "courses.txt"
```

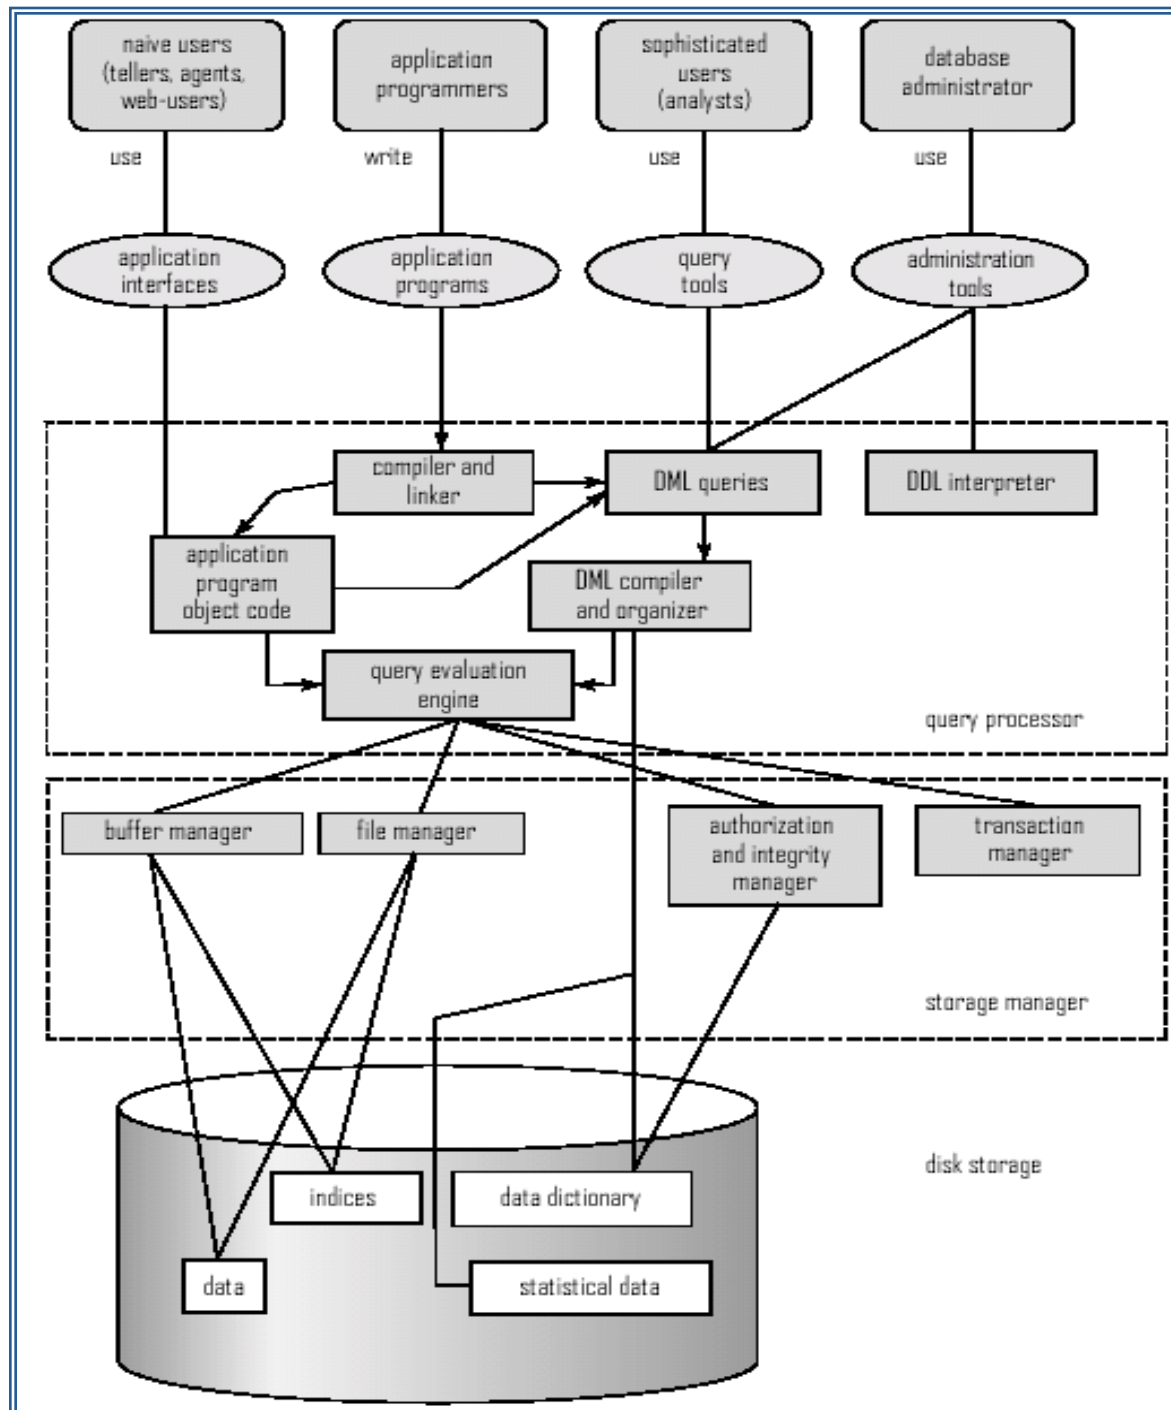
CRASH !

- What is the problem ?
- Large data sets (say 50TB)
 - What is the problem ?
- Simultaneous access by many users
 - Need locks: we know them from OS, but now data on disk; and is there any fun to re-implement them ?
- Data- and user-specific access control

Enter a DBMS

“Two-tier database system”





Main Advantages of DBMS

- Data independence
- Efficient querying
- Data integrity and security
- Data administration
- Concurrency and crash recovery
- Reduced developer bandwidth

Functionality of a DBMS

The programmer sees SQL, which has two components:

- Data Definition Language - DDL
- Data Manipulation Language - DML
 - query language

Behind the scenes the DBMS has:

- Query optimizer
- Query engine
- Storage management
- Transaction Management (concurrency, recovery)

Functionality of a DBMS

Two things to remember:

- Client-server architecture
 - Slow, cumbersome connection
 - But good for the data
 - But the DBMS is *general* and *convenient*
 - We can do any very specific task faster outside the DBMS

How the User Sees the DBMS

- Start with DDL to *create tables*:

```
CREATE TABLE Students (  
    Name CHAR(30)  
    SSN CHAR(9) PRIMARY KEY NOT NULL,  
    Category CHAR(20)  
) ...
```

- Continue with DML to *populate tables*:

```
INSERT INTO Students  
VALUES('Charles', '123456789', 'undergraduate')  
. . . .
```

How the User Sees the DBMS

- Tables:

Students:

SSN	Name	Category
123-45-6789	Charles	undergrad
234-56-7890	Dan	grad

Takes:

SSN	CID
123-45-6789	CSE444
123-45-6789	CSE444
234-56-7890	CSE142
	...

Courses:

CID	Name	Quarter
CSE444	Databases	fall
CSE541	Operating systems	winter

- Still implemented as files, but behind the scenes can be quite complex

“data independence” = separate logical view from physical implementation

Another context...

Accounts

<i>accountNo</i>	<i>balance</i>	<i>type</i>
12345	1000.00	savings
67890	2846.92	checking

```
SELECT balance
FROM Accounts
WHERE accountNo = 67890;
```

```
SELECT accountNo
FROM Accounts
WHERE type = 'savings'
      AND balance < 0;
```

Key observation:

- Regardless of the context:
 - university settings
 - banking settings

The “structure” of interacting is the same...

Transactions

- Enroll “Mary Johnson” in “CSE444”:

```
BEGIN TRANSACTION;  
  
INSERT INTO Takes  
  SELECT Students.SSN, Courses.CID  
  FROM Students, Courses  
  WHERE Students.name = 'Mary Johnson' and  
         Courses.name = 'CSE444'  
  
-- More updates here....  
  
IF everything-went-OK  
  THEN COMMIT;  
ELSE ROLLBACK
```

If system crashes, the transaction is still either committed or aborted

Transactions

- A *transaction* = sequence of statements that either all succeed, or all fail
- Transactions have the ACID properties:
 - A = atomicity
 - C = consistency
 - I = Isolation
 - D = durability

Queries


- Find all courses that “Mary” takes

```
SELECT C.name  
FROM   Students S, Takes T, Courses C  
WHERE  S.name="Mary" and  
        S.ssn = T.ssn and T.cid = C.cid
```

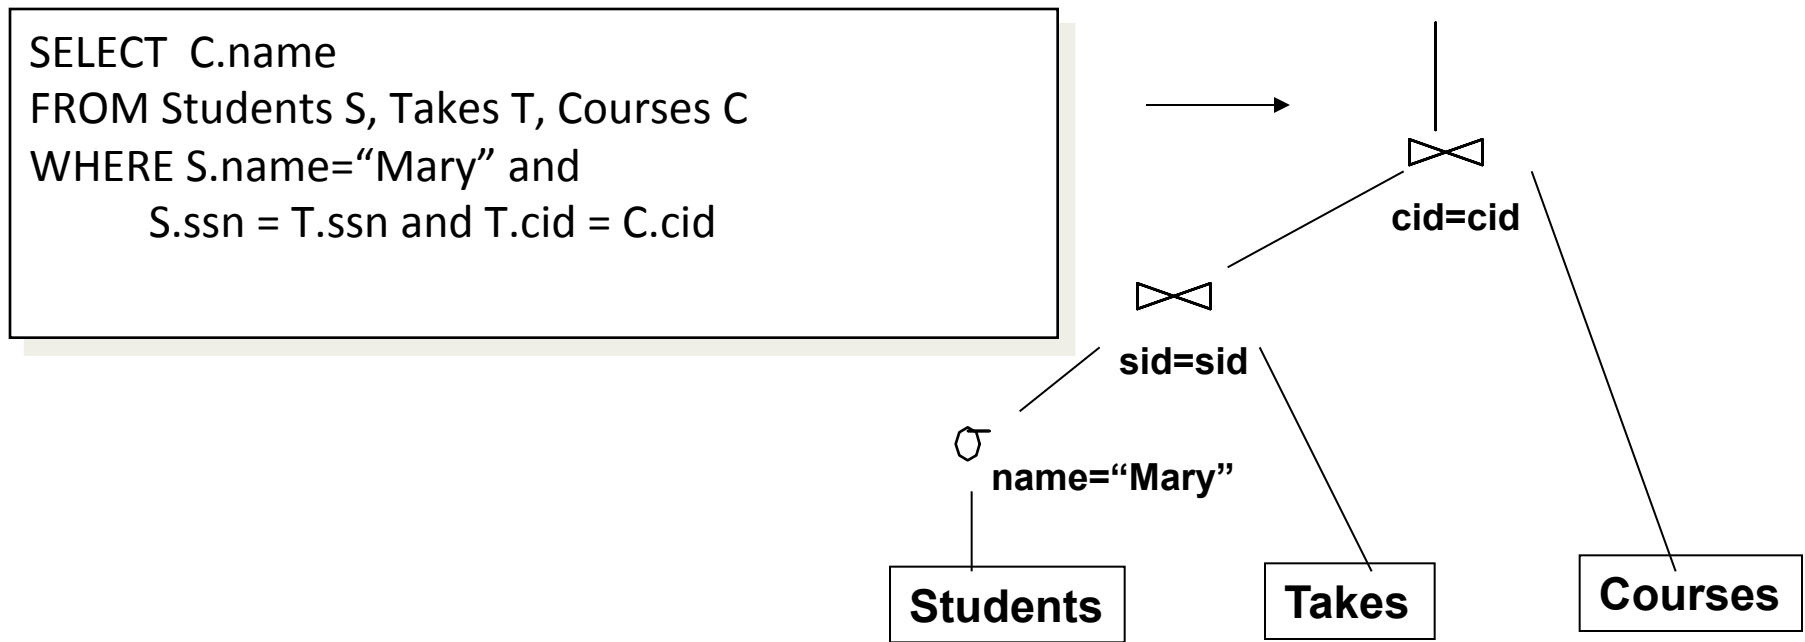
- What happens behind the scene ?
 - Query processor figures out how to answer the query efficiently.

Queries, behind the scene

Declarative SQL query



Imperative query execution plan:



The **optimizer** chooses the best execution plan for a query

People Who Interact With Databases

- Application developers
- Database administrators
- Database producers

Alternative Data Models

- Hierarchical
- Object-oriented
- Object-relational
- Semistructured
- Unstructured
- Multidimensional
- Graphs
- ...many more

Conclusions

- DBMSs are everywhere!
- SQL is the lingua franca of data management
 - But this is slowly changing
- Databases enable people to organize and retrieve their data in a general and consistent fashion.