

Advanced SQL

EECS 339

Lecture 7

The story so far

```
SELECT    [DISTINCT] target-list  
FROM      relation-list  
WHERE     qualification  
GROUP BY  grouping-list  
HAVING    group-qualification
```

- Select-project- (simple) join
- Aggregation

Example Instances

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

- We will use these instances of the Sailors and Reserves relations in our examples.

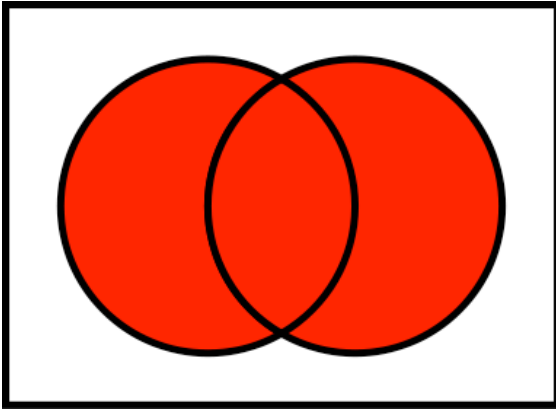
S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

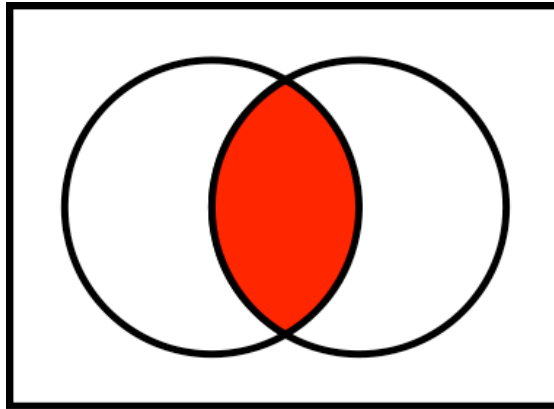
S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

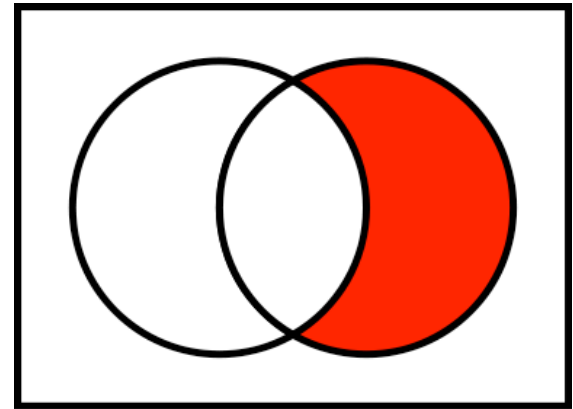
Set Comparisons



UNION



INTERSECT



EXCEPT

Find sid' s of sailors who' ve reserved a red or a green boat

- **UNION**: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).
- If we replace **OR** by **AND** in the first version, what do we get?
- Also available: **EXCEPT** (What do we get if we replace **UNION** by **EXCEPT**?)

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND (B.color= 'red' OR B.color= 'green' )
```

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND B.color= 'red'
```


UNION

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND B.color= 'green'
```

Find sid's of sailors who've reserved a red and a green boat

```
SELECT S.sid
FROM Sailors S, Boats B1, Reserves R1,
     Boats B2, Reserves R2
WHERE S.sid=R1.sid AND R1.bid=B1.bid
AND S.sid=R2.sid AND R2.bid=B2.bid
AND (B1.color='red' AND B2.color='green')
```

- **INTERSECT**: Can be used to compute the intersection of any two *union-compatible* sets of tuples.

```
SELECT S.sid  Key field!
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
AND B.color='red'
```

```
INTERSECT
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
AND B.color='green'
```

- Contrast symmetry of the UNION and INTERSECT queries with how much the other versions differ.

Nested Queries

Find names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                FROM Reserves R
                WHERE R.bid=103)
```

- A very powerful feature of SQL: a WHERE clause can itself contain an SQL query! (Actually, so can FROM and HAVING clauses.)
- To find sailors who've *not* reserved #103, use NOT IN.
- To understand semantics of nested queries, think of a nested loops evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*

Nested Queries with Correlation

Find names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT *
               FROM Reserves R
               WHERE R.bid=103 AND S.sid=R.sid)
```



- **EXISTS** is another set comparison operator, like **IN**.
- If **UNIQUE** is used, and * is replaced by *R.bid*, finds sailors with at most one reservation for boat #103. (**UNIQUE** checks for duplicate tuples; * denotes all attributes. Why do we have to replace * by *R.bid*?)
- Illustrates why, in general, subquery must be re-computed for each Sailors tuple.

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS, UNION ALL and NOT UNIQUE.
- Also available: *op* ANY, *op* ALL, *op* IN
- Find sailors whose rating is greater than that of some sailor called Horatio:

```
SELECT *  
FROM Sailors S  
WHERE S.rating > ANY (SELECT S2.rating  
                      FROM Sailors S2  
                      WHERE S2.sname= 'Horatio')
```

Rewriting INTERSECT Queries Using IN

Find sid's of sailors who've reserved both a red and a green boat:

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color= 'red'
      AND S.sid IN (SELECT S2.sid
                     FROM Sailors S2, Boats B2, Reserves R2
                     WHERE S2.sid=R2.sid AND R2.bid=B2.bid
                        AND B2.color= 'green' )
```

- Similarly, EXCEPT queries re-written using NOT IN.
- To find *names* (not *sid's*) of Sailors who've reserved both red and green boats, just replace *S.sid* by *S.sname* in SELECT clause.

Division in SQL

(1)

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
      ((SELECT B.bid
        FROM Boats B)
      EXCEPT
      (SELECT R.bid
        FROM Reserves R
        WHERE R.sid=S.sid))
```

Find sailors who've reserved all boats.

- Let's do it the hard way, without EXCEPT:

(2) SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS (SELECT B.bid
FROM Boats B

Sailors S such that ...

WHERE NOT EXISTS (SELECT R.bid
FROM Reserves R
WHERE R.bid=B.bid

there is no boat B without ...

AND R.sid=S.sid))

a Reserves tuple showing S reserved B

Top-K Querying

- Look at the first k results from a query
- Use **LIMIT** clause
- Example: Select the top 10 sailors with the largest number of reservations

```
SELECT sid  
FROM sailors s, reservations r  
WHERE s.sid = r.rid  
GROUP BY sid  
ORDER BY COUNT(*)  
LIMIT 10;
```

Study Break: Advanced SQL

- Consider the following schema:
 - Suppliers(sid, sname, address)
 - Parts(pid, color)
 - Catalog(sid, pid, cost)
- Express the following as SQL queries:
 - Find the sids of suppliers who supply a red part and a green part.
 - For each part, find the sname of the supplier who charges the most for that part.
 - For every supplier that supplies a green part and a red part, print the name and price of the most expensive part that she supplies

Summary

- SQL permits us to reason about data a tuple at a time
- We also work with data using set comparisons
- Combine the two to express a rich set of query semantics