

Advanced SQL 2

EECS 339

Lecture 8

Overview

- SQL
 - Join types
 - 3-valued logic
 - Strings and pattern matching
- Query execution
 - JDBC
 - Explain


Join Types

- Joins needed to combine data from multiple tables and expressions
- Recall that join output concatenates the schemas of inputs


Joins

- Equi-join
 - Join on a specific attribute or set of attributes
- Natural join
 - Compare all common columns between 2 relations, select full equalities
- Theta join
 - Allows arbitrary comparisons between attributes, e.g., $<$, \leq , etc.


SQL JOINS




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

Coalesce

- If there is a null, e.g., left join, sometimes want to replace it with a default value
- Do this with coalesce
- E.g., “find the credit card balance of all of a bank’s account holders. If no cc record exists, their debt is zero”.

```
SELECT acct_id, COALESCE(c.balance, 0)
FROM acct_holders LEFT JOIN
credit_cards c ON acct_id;
```

NULL Values

- Tuples in SQL relations can have NULL as a value for one or more components.
- Meaning depends on context. Two common cases:
 - *Missing value* : e.g., we know Joe's Bar has some address, but we don't know what it is.
 - *Inapplicable* : e.g., the value of attribute *spouse* for an unmarried person.

Comparing NULLs to Values

- The logic of conditions in SQL is really 3-valued logic: TRUE, FALSE, UNKNOWN.
- Comparing any value (including NULL itself) with NULL yields UNKNOWN.
- A tuple is in a query answer iff the WHERE clause is TRUE (not FALSE or UNKNOWN).

String Matching

- A condition can compare a string to a pattern by:
 - <Attribute> LIKE <pattern> or <Attribute> NOT LIKE <pattern>
- *Pattern* is a quoted string with % = “any string”; _ = “any character.”

LIKE

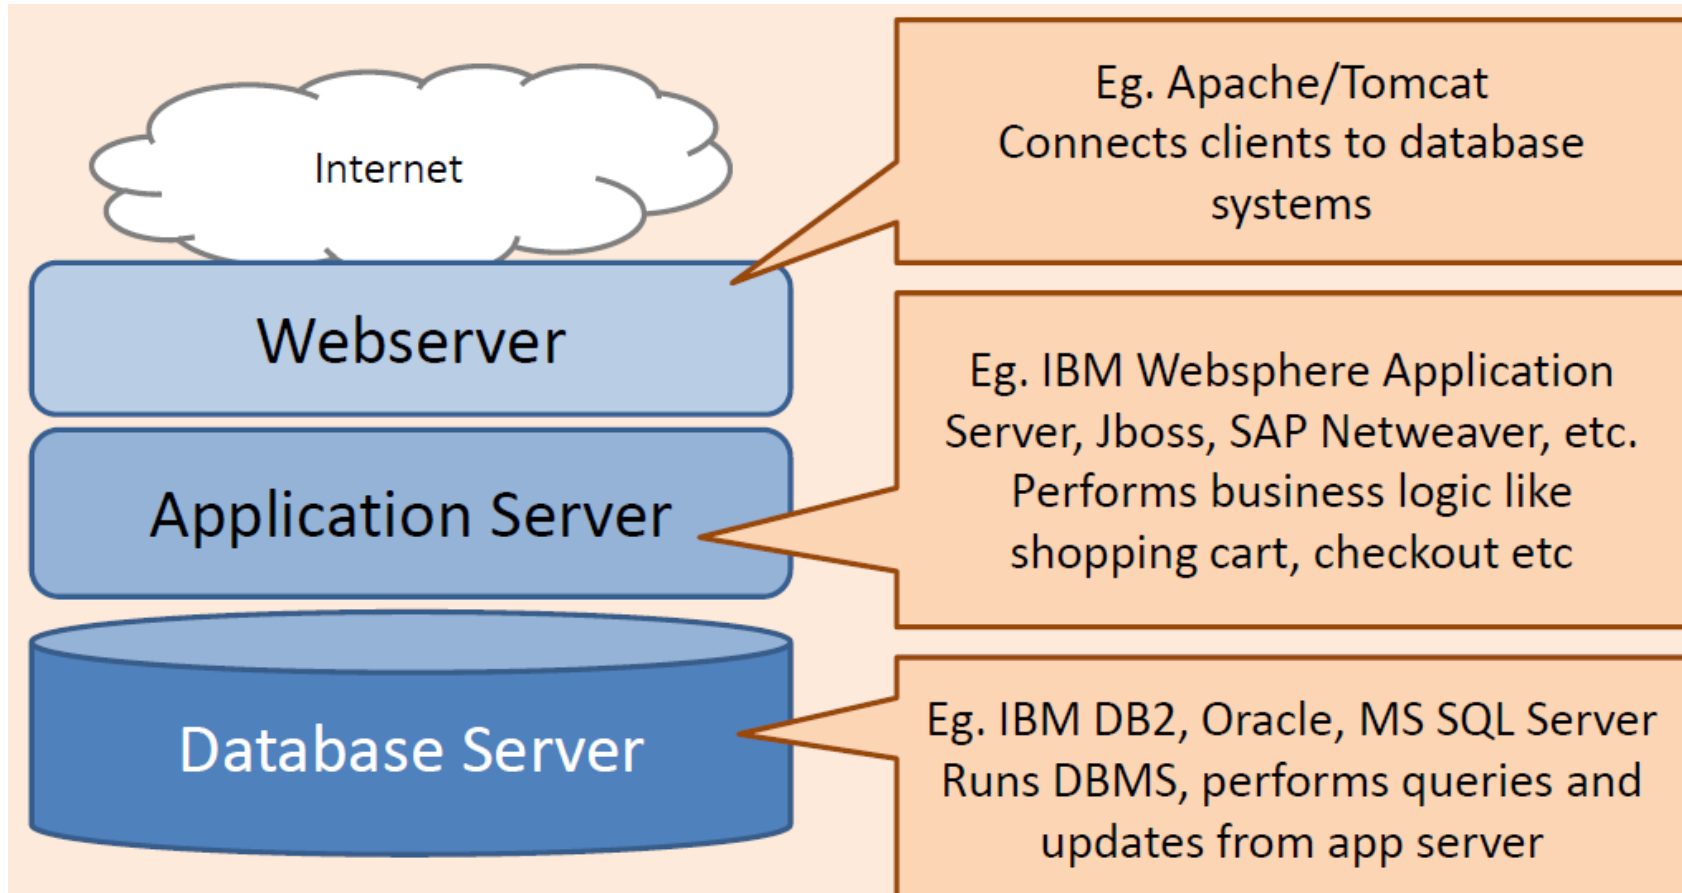
- Using Drinkers(name, addr, phone) find the drinkers with exchange 555:

```
SELECT name  
FROM Drinkers  
WHERE phone LIKE '%555-__ __ __';
```

SQL in Real Programs


- We have seen only how SQL is used at the generic query interface --- an environment where we sit at a terminal and ask queries of a database.
- Reality is almost always different: conventional programs interacting with SQL.

Three Tier Architecture



Large, Internet based enterprises

Basic SQL Environment view...

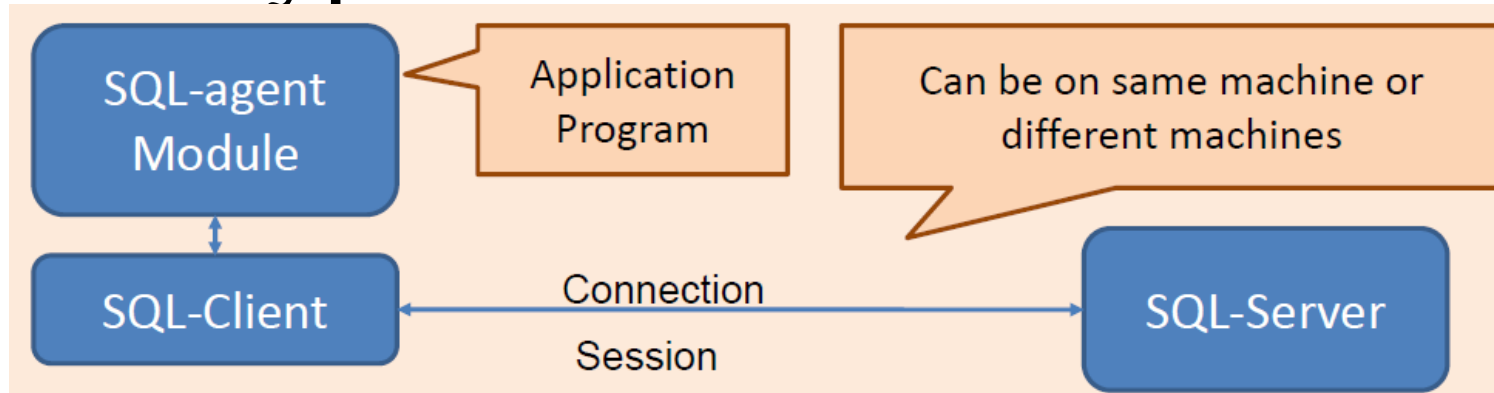


Schemas : tables, views, assertions, triggers

Catalogs : collection of schemas
(aka “databases” in DB2)

Clusters : collection of catalogs (aka “database instance” in DB2)

Typical Client-Server Model



- CONNECT TO <server> AS
<connection name>
AUTHORIZATION

- DISCONNECT/CONNECT
RESET/TERMINATE

- Session – SQL
operations performed
while a connection is
active

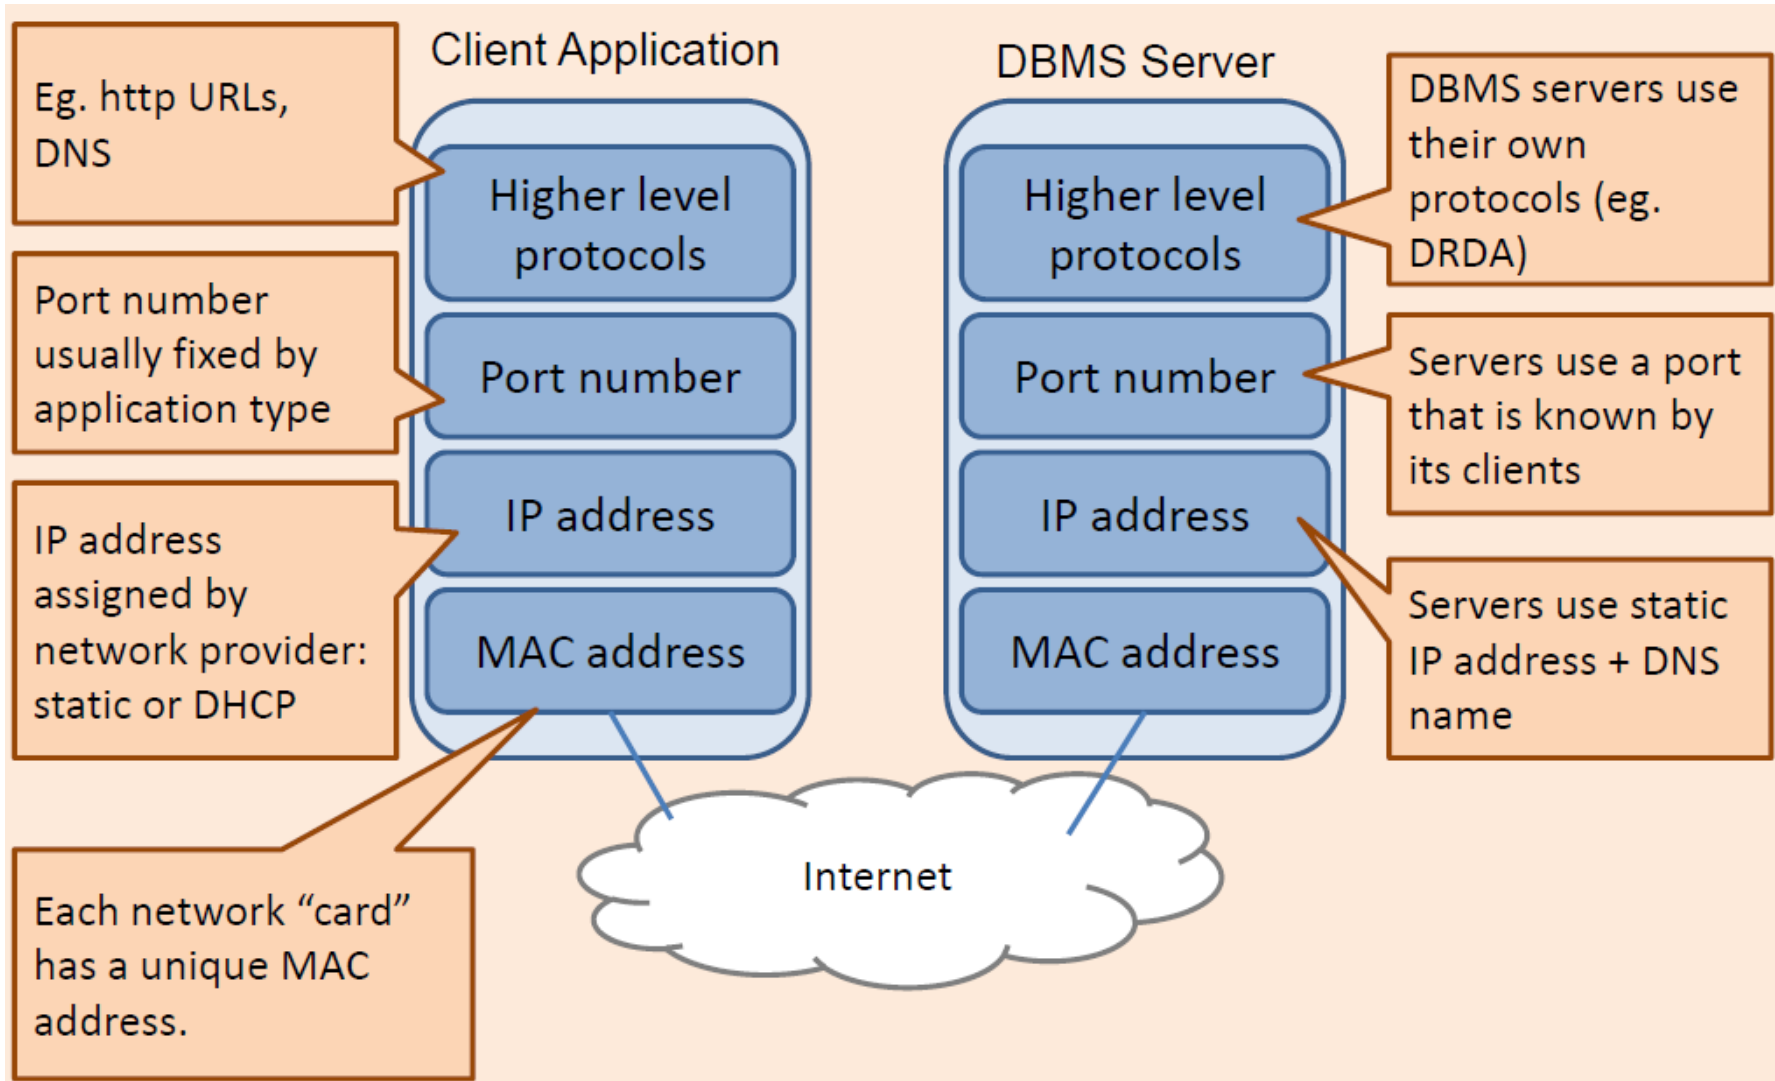
- Programming API:

- Generic SQL Interface
- Embedded SQL in a host language
- True Modules. Eg. Stored procedures.

EXTREMES OF THE INTEGRATION spectrum:

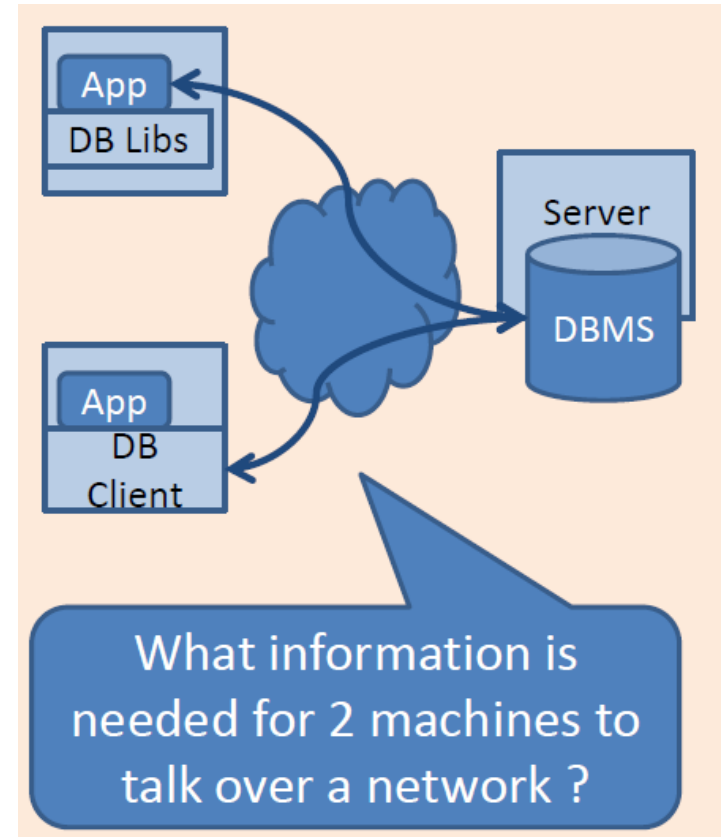
- Highly integrated (e.g. Microsoft linq)
 - Compiler checking of database operations
- Loosely integrated eg. ODBC & JDBC
 - Provides a way to call SQL from host language
 - Host language compiler doesn't understand database operations.
- Requirements:
- *Perform DB operations from host language*
- *DB operations need to access variables in host language*

In context of Networking



Remote Client Access

- Applications run on a machine that is separate from the DB server
- DBMS “thin” client
 - Libraries to which you link your app
 - App needs to know how to talk to DBMS server via network
- DBMS “full” client layer
 - Need to pre-configure the thick client layer to talk to DBMS server
 - Your app talks to a DBMS client layer as if it is talking to the server



Options

1. Code in a specialized language is stored in the database itself (e.g., PSM, PL/SQL).
2. SQL statements are embedded in a *host language* (e.g., C).
3. Connection tools are used to allow a conventional language to access a database (e.g., CLI, JDBC, PHP/DB).

Host/SQL Interfaces Via Libraries

- The third approach to connecting databases to conventional languages is to use library calls.
 1. C + CLI
 2. Java + JDBC
 3. PHP + PEAR/DB

Three-Tier Architecture

- A common environment for using a database has three tiers of processors:
 1. *Web servers* --- talk to the user.
 2. *Application servers* --- execute the business logic.
 3. *Database servers* --- get what the app servers need from the database.


Example: Amazon

- Database holds the information about products, customers, etc.
- Business logic includes things like “what do I do after someone clicks ‘checkout’?”
 - Answer: Show the “how will you pay for this?” screen.

Environments, Connections, Queries

- The database is, in many DB-access languages, an *environment*.
- Database servers maintain some number of *connections*, so app servers can ask queries or perform modifications.
- The app server issues *statements* : queries and modifications, usually.

Diagram to Remember



SQL/CLI

- Instead of using a preprocessor (as in embedded SQL), we can use a library of functions.
 - The library for C is called SQL/CLI = “*Call-Level Interface*.”
 - Embedded SQL’s preprocessor will translate the EXEC SQL ... statements into CLI or similar calls, anyway.

Data Structures

- C connects to the database by structs of the following types:
 - 1. Environments* : represent the DBMS installation.
 - 2. Connections* : logins to the database.
 - 3. Statements* : SQL statements to be passed to a connection.
 - 4. Descriptions* : records about tuples from a query, or parameters of a statement.

Handles

- Function **SQLAllocHandle(T,I,O)** is used to create these structs, which are called environment, connection, and statement *handles*.
 - T = type, e.g., SQL_HANDLE_STMT.
 - I = input handle = struct at next higher level (statement < connection < environment).
 - O = (address of) output handle.

Example: SQLAllocHandle

```
SQLAllocHandle (SQL_HANDLE_STMT,  
               myCon, &myStat);
```

- myCon is a previously created connection handle.
- myStat is the name of the statement handle that will be created.

Preparing and Executing

- **SQLPrepare(H, S, L)** causes the string S , of length L , to be interpreted as a SQL statement and optimized; the executable statement is placed in statement handle H .
- **SQLExecute(H)** causes the SQL statement represented by statement handle H to be executed.

Example: Prepare and Execute

```
SQLPrepare(myStat, "SELECT beer,  
price FROM Sells  
WHERE bar = 'Joe''s Bar'",  
SQL_NTS);  
SQLExecute(myStat);
```

This constant says the second argument is a “null-terminated string”; i.e., figure out the length by counting characters.

Direct Execution

- If we shall execute a statement S only once, we can combine PREPARE and EXECUTE with:

SQLExecuteDirect(H, S, L);

- As before, H is a statement handle and L is the length of string S .

Fetching Tuples

- When the SQL statement executed is a query, we need to fetch the tuples of the result.
 - A cursor is implied by the fact we executed a query; the cursor need not be declared.
- **SQLFetch(*H*)** gets the next tuple from the result of the statement with handle *H*.

Accessing Query Results

- When we fetch a tuple, we need to put the components somewhere.
- Each component is bound to a variable by the function **SQLBindCol**.
 - This function has 6 arguments, of which we shall show only 1, 2, and 4:
 - 1 = handle of the query statement.
 - 2 = column number.
 - 4 = address of the variable.

Example: Binding

- Suppose we have just done **SQLExecute(myStat)**, where myStat is the handle for query

```
SELECT beer, price FROM Sells  
WHERE bar = 'Joe''s Bar'
```

- Bind the result to theBeer and thePrice:

```
SQLBindCol(myStat, 1, , &theBeer, , );
```

```
SQLBindCol(myStat, 2, , &thePrice, , );
```

Example: Fetching

- Now, we can fetch all the tuples of the answer by:

```
while ( SQLFetch(myStat) != SQL_NO_DATA)
{
    /* do something with theBeer and
       thePrice */
}
```

CLI macro representing
SQLSTATE = 02000 = “failed
to find a tuple.”

JDBC

- *Java Database Connectivity* (JDBC) is a library similar to SQL/CLI, but with Java as the host language.
- Like CLI, but with a few notable differences, to be mentioned...

Making a Connection

```
import java.sql.*;  
Class.forName("com.mysql.jdbc.Driver");  
Connection myCon =  
    DriverManager.getConnection(...);
```

The diagram shows four code snippets highlighted with colored boxes: `java.sql.*` (blue), `com.mysql.jdbc.Driver` (pink), `DriverManager` (blue), and `getConnection(...)` (yellow). Arrows point from these boxes to explanatory text: `java.sql.*` points to 'The JDBC classes'; `com.mysql.jdbc.Driver` points to 'The driver for mySql; others exist'; `DriverManager` points to 'Loaded by forName'; and `getConnection(...)` points to 'URL of the database your name, and password go here.'

The JDBC classes

Loaded by
forName

URL of the database
your name, and password
go here.

The driver
for mySql;
others exist

Statements

- JDBC provides two classes:
 - 1. Statement* = an object that can accept a string that is a SQL statement and can execute such a string.
 - 2. PreparedStatement* = an object that has an associated SQL statement ready to execute.

Creating Statements

- The Connection class has methods to create Statements and PreparedStatement.

```
Statement stat1 = myCon.createStatement();
```

```
PreparedStatement stat2 =  
myCon.createStatement(  
    "SELECT beer, price FROM Sells " +  
    "WHERE bar = 'Joe' 's Bar' "  
);
```

`createStatement` with no argument returns a Statement; with one argument it returns a PreparedStatement.

Executing SQL Statements

- JDBC distinguishes queries from modifications, which it calls “updates.”
- Statement and PreparedStatement each have methods `executeQuery` and `executeUpdate`.
 - For Statements: one argument: the query or modification to be executed.
 - For PreparedStatements: no argument.

Example: Update

- stat1 is a Statement.
- We can use it to insert a tuple as:

```
stat1.executeUpdate(  
    "INSERT INTO Sells " +  
    "VALUES ('Brass Rail', 'Bud', 3.00) "  
);
```


Example: Query

- stat2 is a PreparedStatement holding the query
"SELECT beer, price FROM Sells WHERE
bar = 'Joe's Bar'".
- **executeQuery** returns an object of class
ResultSet – we'll examine it later.
- The query:

```
ResultSet menu = stat2.executeQuery();
```

Accessing the ResultSet

- An object of type ResultSet is something like a cursor.
- Method `next()` advances the “cursor” to the next tuple.
 - The first time `next()` is applied, it gets the first tuple.
 - If there are no more tuples, `next()` returns the value `false`.

Accessing Components of Tuples

- When a ResultSet is referring to a tuple, we can get the components of that tuple by applying certain methods to the ResultSet.
- Method `getX(i)`, where X is some type, and i is the component number, returns the value of that component.
 - The value must have type X .

Example: Accessing Components

- Menu = ResultSet for query
“SELECT beer, price FROM Sells WHERE bar = 'Joe'
's Bar' ”.
- Access beer and price from each tuple by:

```
while ( menu.next() ) {  
    theBeer = Menu.getString(1);  
    thePrice = Menu.getFloat(2);  
    /*something with theBeer and  
    thePrice*/  
}
```

Debugging Query Performance

- Sometimes a query does not perform as we expect
- Debug it without running by putting EXPLAIN in front of statement
- Debug it with runtime statistics using EXPLAIN ANALYZE

Conclusions

- SQL has advanced functionality for handling joins, nulls, and strings
- Connection tools are useful for embedded database calls into a program
- EXPLAIN enables us to debug query performance issues