

# Query Evaluation

EECS 339

Lecture 13

# Overview of Query Evaluation

- Plan: *Tree of R.A. ops, with choice of alg for each op.*
  - Each operator typically implemented using a ‘pull’ interface: when an operator is ‘pulled’ for the next output tuples, it ‘pulls’ on its inputs and computes them.
- Two main issues in query optimization:
  - For a given query, **what plans are considered?**
    - Algorithm to search plan space for cheapest (estimated) plan.
  - How is the **cost of a plan estimated?**
- **Ideally**: Want to find best plan. **Practically**: Avoid worst plans!
- We will study the System R approach.

# Some Common Techniques

- Algorithms for evaluating relational operators use some simple ideas extensively:
  - **Indexing:** Can use WHERE conditions to retrieve small set of tuples (selections, joins)
  - **Iteration:** Sometimes, faster to scan all tuples even if there is an index. (And sometimes, we can scan the data entries in an index instead of the table itself.)
  - **Partitioning:** By using sorting or hashing, we can partition the input tuples and replace an expensive operation by similar operations on smaller inputs.

*\* Watch for these techniques as we discuss query evaluation!*

# Statistics and Catalogs

- Need information about the relations and indexes involved. *Catalogs* typically contain at least:
  - # tuples (NTuples) and # pages (NPages) for each relation.
  - # distinct key values (NKeys) and NPages for each index.
  - Index height, low/high key values (Low/High) for each tree index.
- Catalogs updated periodically.
  - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok.
- More detailed information (e.g., histograms of the values in some field) are sometimes stored.

# Access Paths

- ❖ An access path is a method of retrieving tuples:
  - *File scan*, or *index* that *matches* a selection (in the query)
- ❖ A tree index matches (a conjunction of) terms that involve only attributes in a *prefix* of the search key.
  - E.g., Tree index on  $\langle a, b, c \rangle$  *matches* the selection  $a=5 \text{ AND } b=3$ , and  $a=5 \text{ AND } b>6$ , but not  $b=3$ .
- ❖ A hash index matches (a conjunction of) terms that has a term *attribute = value* for every attribute in the search key of the index.
  - E.g., Hash index on  $\langle a, b, c \rangle$  *matches*  $a=5 \text{ AND } b=3 \text{ AND } c=5$ ; but it does not match  $b=3$ , or  $a=5 \text{ AND } b=3$ , or  $a>5 \text{ AND } b=3 \text{ AND } c=5$ .

# A Note on Complex Selections

*(day<8/9/94 AND rname= 'Paul') OR bid=5 OR sid=3*

- Selection conditions are first converted to conjunctive normal form (CNF):  
*(day<8/9/94 OR bid=5 OR sid=3 ) AND  
(rname= 'Paul' OR bid=5 OR sid=3)*
- We only discuss case with no ORs; see text if you are curious about the general case.

# One Approach to Selections

- Find the *most selective access path*, retrieve tuples using it, and apply any remaining terms that don't *match* the index:
  - *Most selective access path*: An index or file scan that we estimate will require the fewest page I/Os.
  - Terms that match this index reduce the number of tuples *retrieved*; other terms are used to discard some retrieved tuples, but do not affect number of tuples/pages fetched.
  - Consider *day<8/9/94 AND bid=5 AND sid=3*. A B+ tree index on *day* can be used; then, *bid=5* and *sid=3* must be checked for each retrieved tuple. Similarly, a hash index on *<bid, sid>* could be used; *day<8/9/94* must then be checked.

# Using an Index for Selections

- Cost depends on #qualifying tuples, and clustering.
  - Cost of finding qualifying data entries (typically small) plus cost of retrieving records (could be large w/o clustering).
  - In example, assuming uniform distribution of names, about 10% of tuples qualify (100 pages, 10000 tuples). With a clustered index, cost is little more than 100 I/Os; if unclustered, up to 10000 I/Os!

```
SELECT *  
FROM   Reserves R  
WHERE  R.rname < 'C%'
```



# Projection

- The expensive part is removing duplicates.
  - SQL systems don't remove duplicates unless the keyword DISTINCT is specified in a query.
- Sorting Approach: Sort on <sid, bid> and remove duplicates. (Can optimize this by dropping unwanted information while sorting.)
- Hashing Approach: Hash on <sid, bid> to create partitions. Load partitions into memory one at a time, build in-memory hash structure, and eliminate duplicates.
- If there is an index with both R.sid and R.bid in the search key, may be cheaper to sort data entries!

```
SELECT  DISTINCT  
        R.sid, R.bid  
FROM    Reserves R
```

# Join: Index Nested Loops

```
foreach tuple r in R do  
  foreach tuple s in S where  $r_i == s_j$  do  
    add <r, s> to result
```

- If there is an index on the join column of one relation (say S), can make it the inner and exploit the index.
  - Cost:  $M + (M * p_R) * \text{cost of finding matching S tuples}$
  - $M = \# \text{pages of R}$ ,  $p_R = \# \text{R tuples per page}$
- For each R tuple, cost of probing S index is about 1.2 for hash index, 2-4 for B+ tree. Cost of then finding S tuples (assuming Alt. (2) or (3) for data entries) depends on clustering.
  - Clustered index: 1 I/O (typical), unclustered: upto 1 I/O per matching S tuple.

# Examples of Index Nested Loops

- Hash-index (Alt. 2) on *sid* of Sailors (as inner):
  - Scan Reserves: 1000 page I/Os,  $100 \times 1000$  tuples.
  - For each Reserves tuple: 1.2 I/Os to get data entry in index, plus 1 I/O to get (the exactly one) matching Sailors tuple.  
Total: 220,000 I/Os.
- Hash-index (Alt. 2) on *sid* of Reserves (as inner):
  - Scan Sailors: 500 page I/Os,  $80 \times 500$  tuples.
  - For each Sailors tuple: 1.2 I/Os to find index page with data entries, plus cost of retrieving matching Reserves tuples.  
Assuming uniform distribution, 2.5 reservations per sailor ( $100,000 / 40,000$ ). Cost of retrieving them is 1 or 2.5 I/Os depending on whether the index is clustered.

# Sophisticated Joins

- Until now, we have used nested loops for joining data
  - This is slow,  $n^2$  comparisons
- How can we do better?
  - Sorting
  - Divide & conquer
- Trade-off in I/O and CPU time for each algo

# Sort Merge Join

Equi-join of two tables S & R

$|S|$  = Pages in S;  $\{S\}$  = Tuples in S

$|S| \geq |R|$

M pages of memory;  $M > \sqrt{|S|}$

Algorithm:

- Partition S and R into memory sized sorted runs, write out to disk
- Merge all runs simultaneously

Total I/O cost: Read  $|R|$  and  $|S|$  twice, write once

**$3(|R| + |S|)$  I/Os**

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4

R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

R1	R2	R3	S1	S2	S3
1 ←	6 ←	1 ←	2 ←	8 ←	4 ←
3	9	7	3	9	6
4	14	11	7	12	15

OUTPUT

Need enough memory to keep 1 page of each run in memory at a time

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

R1	R2	R3	S1	S2	S3
1	6 	1	2 	8 	4 
3 	9	7 	3	9	6
4	14	11	7	12	15

OUTPUT

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

R1	R2	R3	S1	S2	S3
1	6 	1	2	8 	4 
3 	9	7 	3 	9	6
4	14	11	7	12	15

OUTPUT



# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

R1	R2	R3	S1	S2	S3
1	6 	1	2	8 	4 
3 	9	7 	3 	9	6
4	14	11	7	12	15

OUTPUT
(3,3)

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

R1	R2	R3	S1	S2	S3
1	6 	1	2	8 	4 
3	9	7 	3 	9	6
4 	14	11	7	12	15

OUTPUT
(3,3)
(4,4)

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

R1	R2	R3	S1	S2	S3
1	6 	1	2	8 	4 
3	9	7 	3	9	6
4 	14	11	7 	12	15

OUTPUT
(3,3)
(4,4)

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

R1	R2	R3	S1	S2	S3
1	6 	1	2	8 	4
3	9	7 	3	9	6 
4 	14	11	7 	12	15

OUTPUT
(3,3)
(4,4)
(6,6)

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

R1	R2	R3	S1	S2	S3
1	6 	1	2	8 	4
3	9	7 	3	9	6 
4 	14	11	7 	12	15

...

OUTPUT
(3,3)
(4,4)
(6,6)
(7,7)

Output in  
sorted  
order!

# Study Break: Sort-Merge Join

- Say we are joining tables:

A=8,20,19,20,3,13,20,18,6,5,4,5

B=15,1,3,13,13,10,19,6,8,15,16,2

If our in-memory runs are of length 4, what will the sorted runs be for A and B?

What is the join output? Walk through the steps of the merge.

# Study Break Solution

- A: (8,19,20,20) (3,13,18,20)(4,5,5,6)
- B: (1,3,13,15)(6,10,13,19)(2,8,15,16)

Output: (3,3)(6,6) (8,8)(13,13)(13,13)(19,19)  
(output in sorted order)

# Simple Hash

Algorithm:

Given hash function  $H(x) \rightarrow [0, \dots, P-1]$

where  $P$  is number of partitions

for  $i$  in  $[0, \dots, P-1]$ :

for each  $r$  in  $R$ :

if  $H(r)=i$ , add  $r$  to in-memory hash

otherwise, write  $r$  back to disk in  $R'$

for each  $s$  in  $S$ :

if  $H(s)=i$ , lookup  $s$  in hash, output matches

otherwise, write  $s$  back to disk in  $S'$

replace  $R$  with  $R'$ ,  $S$  with  $S'$



# Simple Hash I/O Analysis

Suppose  $P=2$ , and hash uniformly maps tuples to partitions

Read  $|R| + |S|$

Write  $1/2 (|R| + |S|)$

Read  $1/2 (|R| + |S|) = 2 (|R| + |S|)$

$P=3$

Read  $|R| + |S|$

Write  $2/3 (|R| + |S|)$

Read  $2/3 (|R| + |S|)$

Write  $1/3 (|R| + |S|)$

Read  $1/3 (|R| + |S|) = 3 (|R| + |S|)$

$P=4$

$$|R| + |S| + 2 * (3/4 (|R| + |S|)) + 2 * (2/4 (|R| + |S|)) + 2 * (1/4 (|R| + |S|)) \\ = 4 (|R| + |S|)$$

➔  $P = n ; n * (|R| + |S|) \text{ I/Os}$

# Grace Hash

Algorithm:

Partition:

Suppose we have  $P$  partitions, and  $H(x) \rightarrow [0 \dots P-1]$

Choose  $P = |S| / M \rightarrow P \leq \sqrt{|S|}$  //may need to leave a little slop for imperfect hashing

Allocate  $P$  1-page output buffers, and  $P$  output files for  $R$

For each  $r$  in  $R$ :

Write  $r$  into buffer  $H(r)$

If buffer full, append to file  $H(r)$

Allocate  $P$  output files for  $S$

For each  $s$  in  $S$ :

Write  $s$  into buffer  $H(s)$

if buffer full, append to file  $H(s)$

Join:

For  $i$  in  $[0, \dots, P-1]$

Read file  $i$  of  $R$ , build hash table

Scan file  $i$  of  $S$ , probing into hash table and outputting matches

Total I/O cost: Read  $|R|$  and  $|S|$  twice, write once

**$3(|R| + |S|)$  I/Os**

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2

P output buffers

F0	F1	F2

P output files

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
		5

F0	F1	F2

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
	4	5

F0	F1	F2

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
3	4	5

F0	F1	F2

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
3	4	5
6		

F0	F1	F2

# Example

$P = 3; H(x) = x \bmod P$



$R = 5, 4, 3, 6, 9, 14, 1, 7, 11$

$S = 2, 3, 7, 12, 9, 8, 4, 15, 6$

R0	R1	R2
3	4	5
6		

Need to flush R0 to F0!

F0	F1	F2



# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
	4	5

F0	F1	F2
3		
6		

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
9	4	5

F0	F1	F2
3		
6		

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
9	4	5
		14

F0	F1	F2
3		
6		

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
9	4	5
	1	14

F0	F1	F2
3		
6		

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
9	4	5
	1	14

F0	F1	F2
3		
6		

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
9		5
		14

F0	F1	F2
3	4	
6	1	

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
9	7	5
		14

F0	F1	F2
3	4	
6	1	

# Example

$P = 3; H(x) = x \bmod P$



$R = 5, 4, 3, 6, 9, 14, 1, 7, 11$

$S = 2, 3, 7, 12, 9, 8, 4, 15, 6$

R0	R1	R2
9	7	5
		14

F0	F1	F2
3	4	
6	1	



# Example

$P = 3; H(x) = x \bmod P$



$R = 5, 4, 3, 6, 9, 14, 1, 7, 11$

$S = 2, 3, 7, 12, 9, 8, 4, 15, 6$

R0	R1	R2
9	7	

F0	F1	F2
3	4	5
6	1	14

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R0	R1	R2
9	7	11

F0	F1	F2
3	4	5
6	1	14

# Example

$P = 3; H(x) = x \bmod P$



$R = 5, 4, 3, 6, 9, 14, 1, 7, 11$

$S = 2, 3, 7, 12, 9, 8, 4, 15, 6$

R0	R1	R2

F0	F1	F2
3	4	5
6	1	14
9	7	11

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11



S=2,3,7,12,9,8,4,15,6

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		

# Example

$$P = 3; H(x) = x \bmod P$$

Matches:

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

Load F0 from R into memory

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		

# Example

$$P = 3; H(x) = x \bmod P$$

Matches:

R=5,4,3,6,9,14,1,7,11


S=2,3,7,12,9,8,4,15,6

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

Load F0 from R into memory

S Files



F0	F1	F2
3	7	2
12	4	8
9		
15		
6		

Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

Matches:  
3,3

R=5,4,3,6,9,14,1,7,11


S=2,3,7,12,9,8,4,15,6

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

Load F0 from R into memory

S Files



F0	F1	F2
3	7	2
12	4	8
9		
15		
6		

Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

Matches:  
3,3

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

Load F0 from R into memory

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		

Scan F0 from S





# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

Load F0 from R into memory

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		



Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

Load F0 from R into memory

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		



Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

6,6

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

Load F0 from R into memory

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		



Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

6,6

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

6,6

7,7

4,4

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

6,6

7,7

4,4

R Files

F0	F1	F2
3	4	5
6	1	14
9	7	11

S Files

F0	F1	F2
3	7	2
12	4	8
9		
15		
6		

# Execution Costs

Notation:  $P$  partitions / passes over data; assuming hash is  $O(1)$

Sort-Merge	Simple Hash	Grace Hash
I/O: $3( R  +  S )$ CPU: $O(P \times \{S\}/P \log \{S\}/P)$	I/O: $P( R  +  S )$ CPU: $O(\{R\} + \{S\})$	I/O: $3( R  +  S )$ CPU: $O(\{R\} + \{S\})$

Grace hash is generally a safe bet, unless memory is close to size of tables, in which case simple can be preferable

Extra cost of sorting makes sort merge unattractive unless there is a way to access tables in sorted order (e.g., a clustered index), or a need to output data in sorted order (e.g., for a subsequent ORDER BY)

# Study Break: Grace Hash Join

- Say we are joining tables:

A=8,20,19,20,3,13,20,18,6,5,4,5

B=15,1,3,13,13,10,19,6,8,15,16,2

If we have four partitions, what will the hash partitions be for A and B?

Walk through the steps for producing this join's output.



# Study Break Solution

- Partitions:
  - A0: (8,20,20,20,4) A1: (13,5,5) A2: (18,6) A3: (19,3)
  - B0: (8,16) B1: (1,13,13) B2: (2, 6, 10) B3: (15,3,19,15)
- Execution:
  - Join A0 w/B0 (produces (8,8))
  - Join A1 w/B1 (produces (13,13), (13,13))
  - ...

# Highlights of System R Optimizer

- Impact:
  - Most widely used currently; works well for  $< 10$  joins.
- **Cost estimation:** Approximate art at best.
  - Statistics, maintained in system catalogs, used to estimate cost of operations and result sizes.
  - Considers combination of CPU and I/O costs.
- **Plan Space:** Too large, must be pruned.
  - Only the space of *left-deep plans* is considered.
    - Left-deep plans allow output of each operator to be pipelined into the next operator without storing it in a temporary relation.
  - Cartesian products avoided.

# Cost Estimation

- For each plan considered, must estimate cost:
  - Must *estimate cost* of each operation in plan tree.
    - Depends on input cardinalities.
    - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins, etc.)
  - Must also *estimate size of result* for each operation in tree!
    - Use information about the input relations.
    - For selections and joins, assume independence of predicates.

# Size Estimation and Reduction Factors

```
SELECT attribute list  
FROM relation list  
WHERE term1 AND ... AND termk
```

- Consider a query block:
- Maximum # tuples in result is the product of the cardinalities of relations in the FROM clause.
- *Reduction factor (RF)* associated with each *term* reflects the impact of the *term* in reducing result size.  
*Result cardinality = Max # tuples \* product of all RF's.*
  - Implicit *assumption* that *terms are independent!*
  - Term *col=value* has RF  $1/NKeys(I)$ , given index *I* on *col*
  - Term *col1=col2* has RF  $1/MAX(NKeys(I1), NKeys(I2))$
  - Term *col>value* has RF  $(High(I)-value)/(High(I)-Low(I))$

# Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)


Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- Similar to old schema; *rname* added for variations.
- Reserves:
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- Sailors:
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

# Motivating Example

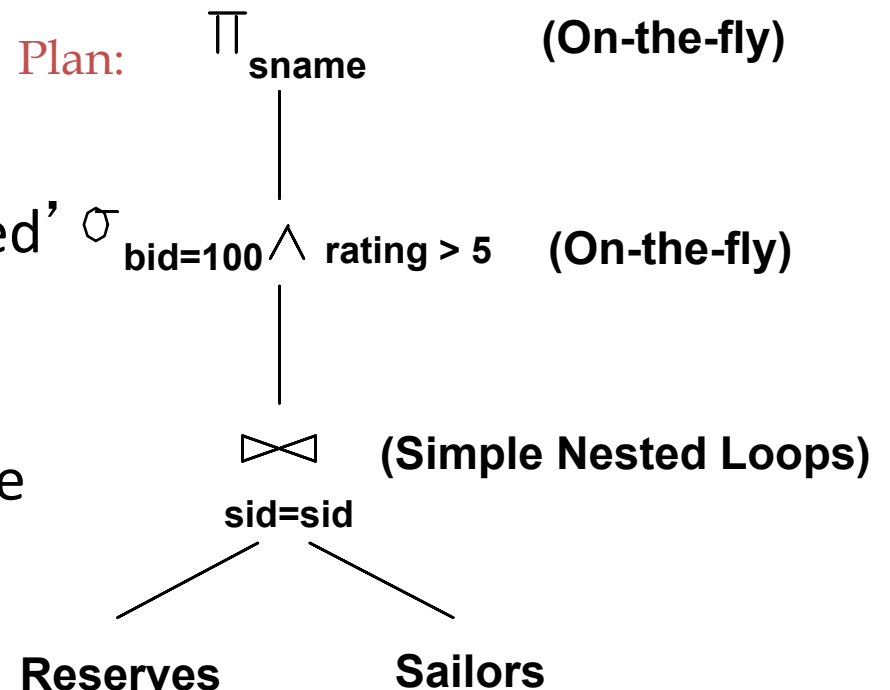
```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

RA Tree:

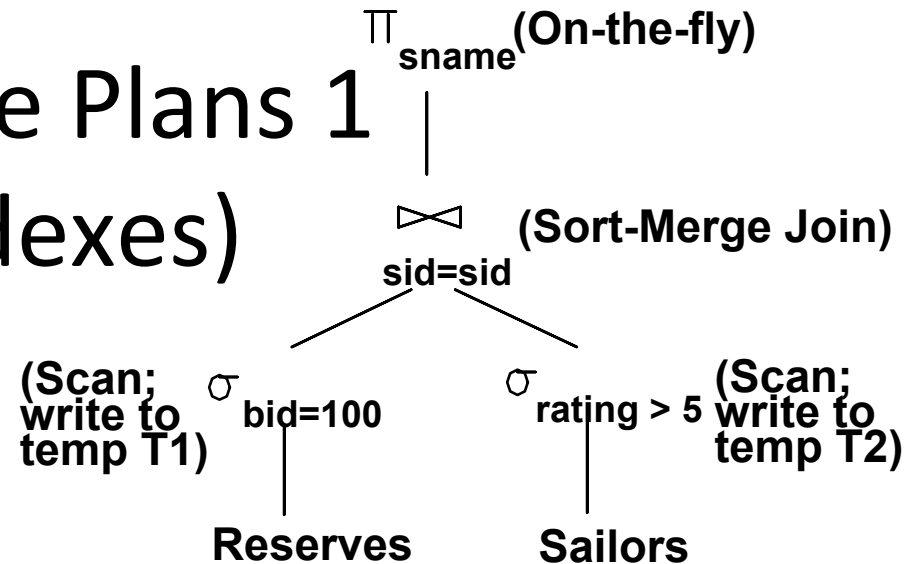


- Cost: 500+500\*1000 I/Os
- By no means the worst plan!
- Misses several opportunities: selections could have been 'pushed' earlier, no use is made of any available indexes, etc.
- *Goal of optimization:* To find more efficient plans that compute the same answer.

Plan:



# Alternative Plans 1 (No Indexes)



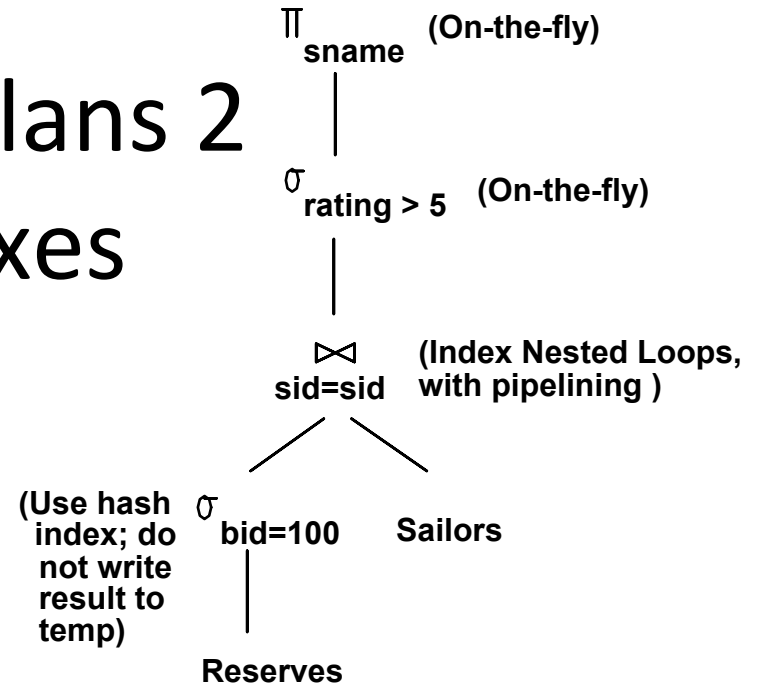
- **Main difference: push selects.**
- With 5 buffers, **cost of plan:**
  - Scan Reserves (1000) + write temp T1 (10 pages, if we have 100 boats, uniform distribution).
  - Scan Sailors (500) + write temp T2 (250 pages, if we have 10 ratings).
  - Sort T1 (2\*2\*10), sort T2 (2\*3\*250), merge (10+250)
  - **Total: 3560 page I/Os.**
- If we used BNL join, join cost = 10+4\*250, **total cost = 2770.**
- If we **'push' projections**, T1 has only *sid*, T2 only *sid* and *sname*:
  - T1 fits in 3 pages, cost of BNL drops to under 250 pages, **total < 2000.**

# Alternative Plans 2 With Indexes

- With clustered index on *bid* of Reserves, we get  $100,000/100 = 1000$  tuples on  $1000/100 = 10$  pages.

- INL with pipelining (outer is not materialized).

–Projecting out unnecessary fields from outer doesn't help.



- ❖ Join column *sid* is a key for Sailors.
  - At most one matching tuple, unclustered index on *sid* OK.
- ❖ Decision not to push *rating>5* before the join is based on availability of *sid* index on Sailors.
- ❖ **Cost:** Selection of Reserves tuples (10 I/Os); for each, must get matching Sailors tuple ( $1000 \times 1.2$ ); total **1210 I/Os**.



# Summary

- There are several alternative evaluation algorithms for each relational operator.
  - Especially for joins!
- A query is evaluated by converting it to a tree of operators and evaluating the operators in the tree.
- Must understand query optimization in order to fully understand the performance impact of a given database design (relations, indexes) on a workload (set of queries).
- Two parts to optimizing a query:
  - Consider a set of alternative plans.
    - Must prune search space; typically, left-deep plans only.
  - Must estimate cost of each plan that is considered.
    - Must estimate size of result and cost for each plan node.
    - *Key issues*: Statistics, indexes, operator implementations.