

Transactions

EECS 339

Lecture 15

Transactions

- Concurrent execution of user programs is essential for good DBMS performance.
 - Because disk accesses are frequent, and relatively slow, it is important to keep the cpu not idle by working on several user programs concurrently.
- A user's program may carry out many operations on the data retrieved from the database, but the DBMS is only concerned about what data is read/written from/to the database.
- A transaction is the DBMS's abstract view of a user program: a sequence of reads and writes.

Transactions

- What's hard?
 - ACID
 - Concurrency control
 - Recovery

Concurrency in a DBMS

- Users submit transactions, and can think of each transaction as executing by itself.
 - Concurrency is achieved by the DBMS, which interleaves actions (reads/writes of DB objects) of various transactions.
 - Each transaction must leave the database in a consistent state if the DB is consistent when the transaction begins.
 - DBMS will enforce some ICs, depending on the ICs declared in CREATE TABLE statements.
 - Beyond this, the DBMS does not really understand the semantics of the data. (e.g., it does not understand how the interest on a bank account is computed).

Atomicity of Transactions

- A transaction might *commit* after completing all its actions, or it could *abort* (or be aborted by the DBMS) after executing some actions.
- A very important property guaranteed by the DBMS for all transactions is that they are *atomic*. That is, a user can think of a Xact as always executing all its actions in one step, or not executing any actions at all.
 - DBMS *logs* all actions so that it can *undo* the actions of aborted transactions.

Example

- Consider two transactions (*Xacts*):

| |
|--|
| T1: BEGIN A=A+100, B=B-100 END |
| T2: BEGIN A=1.06*A, B=1.06*B END |

- ❖ Intuitively, the first transaction is transferring \$100 from B' s account to A' s account. The second is crediting both accounts with a 6% interest payment.
- ❖ There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together. However, the net effect *must* be equivalent to these two transactions running serially in some order.

Example (Cont'd)

- Consider a possible interleaving (schedule):

| | | |
|-----|-------------|------------|
| T1: | $A=A+100,$ | $B=B-100$ |
| T2: | $A=1.06*A,$ | $B=1.06*B$ |

- ❖ This is OK. But what about:

| | | |
|-----|----------------------|-----------|
| T1: | $A=A+100,$ | $B=B-100$ |
| T2: | $A=1.06*A, B=1.06*B$ | |

- ❖ The DBMS' s view of the second schedule:

| | | |
|-----|--------------------------|--------------|
| T1: | $R(A), W(A),$ | $R(B), W(B)$ |
| T2: | $R(A), W(A), R(B), W(B)$ | |

Anomalies with Interleaved Execution

- Reading Uncommitted Data (WR Conflicts, “dirty reads”):

| | | |
|-----|--------------------|-------------------|
| T1: | R(A), W(A), | R(B), W(B), Abort |
| T2: | R(A), W(A), Commit | |

- Unrepeatable Reads (RW Conflicts):

| | | |
|-----|--------------------|--------------------|
| T1: | R(A), | R(A), W(A), Commit |
| T2: | R(A), W(A), Commit | |

Anomalies (Continued)

- Overwriting Uncommitted Data (WW Conflicts):

| | |
|-----------|--------------|
| T1: W(A), | W(B), Commit |
| T2: W(A), | W(B), Commit |

View Serializability

A particular ordering of instructions in a schedule S is *view equivalent* to a serial ordering S' iff:

- Every value read in S is the same value that was read by the same read in S' .
- The final write of every object is done by the same transaction T in S and S'

Conflict Serializability

A schedule is *conflict serializable* if it is possible to swap non-conflicting operations to derive a serial schedule.

For all pairs of conflicting operations {O1 in T1, O2 in T2} either

- O1 always precedes O2, or
- O2 always precedes O1.

Precedence Graph

Given transactions T_i and T_j ,
Create an edge from $T_i \rightarrow T_j$ if:

- T_i reads/writes some A before T_j writes A , or
- T_i writes some A before T_j reads A

If there are cycles in this graph, schedule is not conflict serializable

Study Break: Serializable Transactions

Part A

| | |
|-----------|-----------|
| <u>T1</u> | <u>T2</u> |
| R2 | |
| | R2 |
| W4 | |
| | W2 |
| W2 | |
| | R1 |
| R1 | |
| | W1 |

Is this schedule conflict serializable?

Draw precedence graph.

Is it view serializable?

Part B

| | | |
|-----------|-----------|-----------|
| <u>T1</u> | <u>T2</u> | <u>T3</u> |
| R1 | | |
| R3 | | |
| | W4 | |
| | W2 | |
| | | R2 |
| | | R1 |
| W1 | | |
| W3 | | |
| | R4 | |
| | R2 | |
| | | W3 |

Study Break Solution

Part A

- Conflict serializable? No, $T1(R2) \rightarrow T2(W2) \rightarrow T1(W2)$, Prec. Graph: $T1(R2) \leftrightarrow T2(R2)$,
- View Serializable? No, schedule $T1T2$ impossible b/c of $T2(R2)$. $(T2, T1)$ not possible b/c $T1(R1)$ before $T2(W1)$

Part B

- Conflict Serializable? No. PG: $T1(W3) \rightarrow T3, T1 \leftarrow (R1) T3$,
- View Serializable? No, $(T1, T2, T3)$ not possible b/c $T3$ $R1$ has the wrong value. $T3, T1$ not possible because $R3$ in $T1$ fails to read the right value

Lock-Based Concurrency Control

- Strict Two-phase Locking (Strict 2PL) Protocol:
 - Each Xact must obtain a *S (shared) lock* on object before reading, and an *X (exclusive) lock* on object before writing.
 - All locks held by a transaction are released when the transaction completes
 - *(Non-strict) 2PL Variant:* Release locks anytime, but cannot acquire locks after releasing any lock.
 - If an Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object.
- Strict 2PL allows only serializable schedules.
 - Additionally, it simplifies transaction aborts
 - *(Non-strict) 2PL* also allows only serializable schedules, but involves more complex abort processing

Aborting a Transaction

- If a transaction T_i is aborted, all its actions have to be undone. Not only that, if T_j reads an object last written by T_i , T_j must be aborted as well!
- Most systems avoid such *cascading aborts* by releasing a transaction's locks only at commit time.
 - If T_i writes an object, T_j can read this only after T_i commits.
- In order to *undo* the actions of an aborted transaction, the DBMS maintains a *log* in which every write is recorded. This mechanism is also used to recover from system crashes: all active Xacts at the time of the crash are aborted when the system comes back up.

The Log

- The following actions are recorded in the log:
 - *Ti writes an object*: the old value and the new value.
 - Log record must go to disk before the changed page!
 - *Ti commits/aborts*: a log record indicating this action.
- Log records are chained together by Xact id, so it's easy to undo a specific Xact.
- Log is often *duplexed* and *archived* on stable storage.
- All log related activities (and in fact, all CC related activities such as lock/unlock, dealing with deadlocks etc.) are handled transparently by the DBMS.

Recovering From a Crash

- There are 3 phases in the *Aries* recovery algorithm:
 - Analysis: Scan the log forward (from the most recent *checkpoint*) to identify all Xacts that were active, and all dirty pages in the buffer pool at the time of the crash.
 - Redo: Redoes all updates to dirty pages in the buffer pool, as needed, to ensure that all logged updates are in fact carried out and written to disk.
 - Undo: The writes of all Xacts that were active at the crash are undone (by restoring the *before value* of the update, which is in the log record for the update), working backwards in the log. (Some care must be taken to handle the case of a crash occurring during the recovery process!)

Summary

- Concurrency control and recovery are among the most important functions provided by a DBMS.
- Users need not worry about concurrency.
 - System automatically inserts lock/unlock requests and schedules actions of different Xacts in such a way as to ensure that the resulting execution is equivalent to executing the Xacts one after the other in some order.
- Write-ahead logging (WAL) is used to undo the actions of aborted transactions and to restore the system to a consistent state after a crash.
 - *Consistent state*: Only the effects of committed Xacts seen.