

Optimistic Concurrency Control

EECS 339

Lecture 17

Basic Two Phase Locking Protocol

- Before every read, acquire a shared lock
- Before every write, acquire an exclusive lock (or "upgrade") a shared to an exclusive lock
- Only release locks after all locks have been acquired
- Lock point = point when all locks are acquired
 - Transaction can run to completion
 - Will be serialized after any conflicting transactions that already completed, and before any other conflict transaction

Motivating Example (2 x RX WX RY WY)

[illegible]

Motivating Example (2 x RX WX RY WY)

[illegible]

Motivating Example (2 x RX WX RY WY)

T1	T2
Lock X	
Read X	

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	
Write X	

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	
Write X	
Release X?	

If T1 releases X at this point, T2 could update X and Y before T1, resulting in a non-serial schedule.

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	
Write X	
Release X? Slock Y	

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	
Write X	
Release X? Slock Y	

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	
Write X	
Release X? Slock Y	
Release X	
	Slock X

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	
Write X	
Release X? Slock Y	
Release X	
	Slock X
	Read X
	Write X

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	
Write X	
Release X? Slock Y	
Release X	
	Slock X
	Read X
	Write X
	Acq Y?

Motivating Example (2 x RX WX RY WY)

T1	T2
Slock X	
Read X	
Xlock X	
Write X	
Release X? Slock Y	
Release X	
	Slock X
	Read X
	Write X
	Acq Y? (Waits for T1)

Two Problems

- Cascading Aborts
- Deadlocks

Rigorous Two Phase Locking Protocol

- Before every read, acquire a shared lock
- Before every write, acquire an exclusive lock (or "upgrade") a shared to an exclusive lock
- Release locks only after the transaction commits
- Ensures cascadeless-ness, and that commit order = serialization order

Optimistic Concurrency Control

- No locks
- Check for conflict at commit time
- Each xaction stores its writes in a private workspace
- Keep track of all objs read/written by a xaction

OCC Write

```
twrite(object,value):
```

```
    if object not in write_set: // never written, make  
copy
```

```
        m = read(object)
```

```
        copies[object] = m
```

```
        write_set = write_set U {object}
```

```
write(copies[object], value)
```

OCC Read

```
tread(object):  
    read_set = read_set U {object};  
    if object in write_set:  
        return read(copies[object]);  
    else:  
        return read(object);
```

Validation Rules

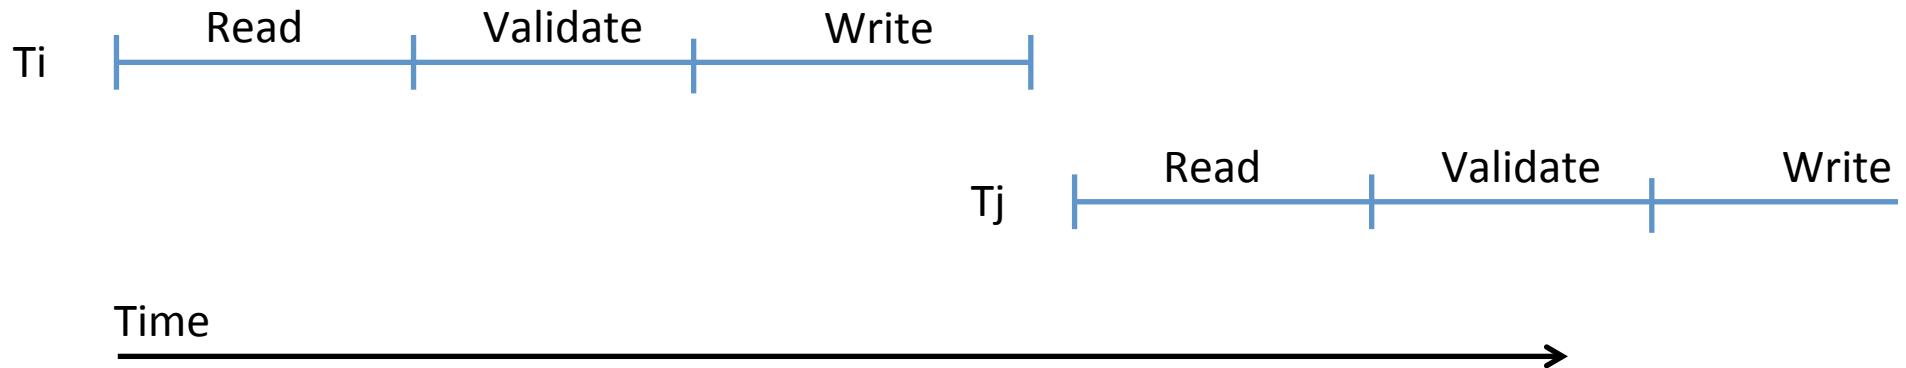
When T_j completes its read phase, require that for all $T_i < T_j$, one of the following conditions must be true for validation to succeed (T_j to commit):

- 1) T_i completes its write phase before T_j starts its read phase
- 2) $W(T_i)$ does not intersect $R(T_j)$, and T_i completes its write phase before T_j starts its write phase.
- 3) $W(T_i)$ does not intersect $R(T_j)$ or $W(T_j)$, and T_i completes its read phase before T_j completes its read phase.
- 4) $W(T_i)$ does not intersect $R(T_j)$ or $W(T_j)$, and $W(T_j)$ does not intersect $R(T_i)$ [no conflicts]

These rules will ensure serializability, with T_j being ordered after T_i with respect to conflicts

Condition 1

T_i completes its write phase before T_j starts its read phase

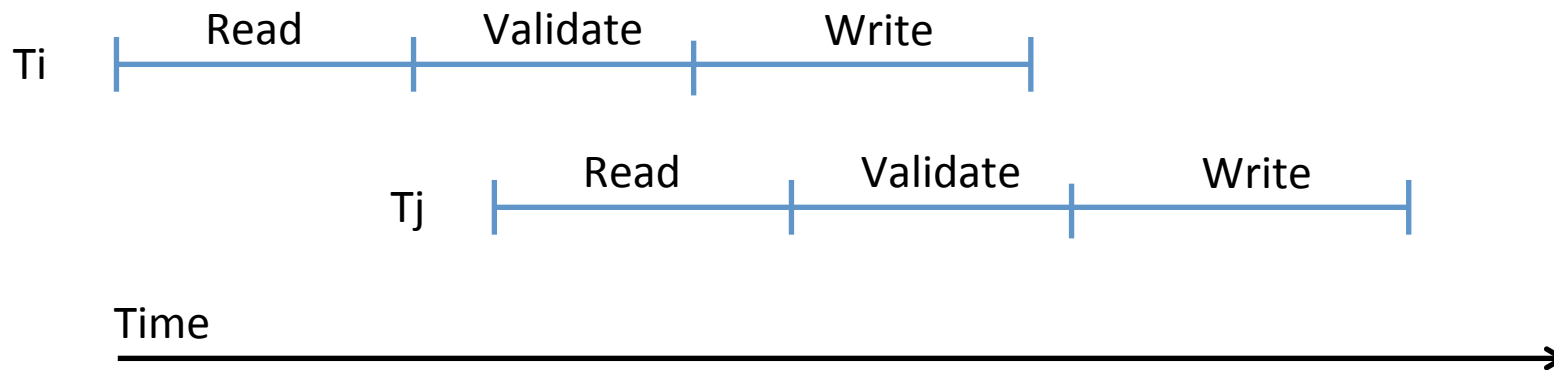


Don't overlap at all.

Condition 2

$W(T_i)$ does not intersect $R(T_j)$, and T_i completes its write phase before T_j starts its write phase.

$$W(T_i) \cap R(T_j) = \{ \} \quad R(T_i) \cap W(T_j) \neq \{ \} \quad W(T_i) \cap W(T_j) \neq \{ \}$$



T_j doesn't read anything T_i wrote.

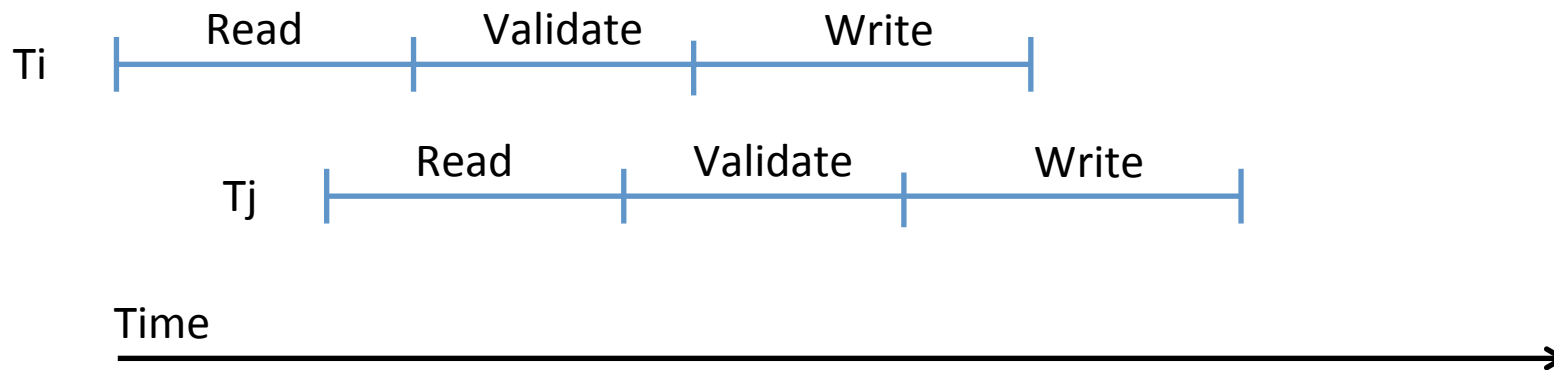
Anything T_j wrote that T_i also wrote will be installed afterwards.

Anything T_i read will not reflect T_j 's writes

Condition 3

$W(T_i)$ does not intersect $R(T_j)$ or $W(T_j)$, and T_i completes its read phase before T_j completes its read phase.

$$W(T_i) \cap R(T_j) = \{ \} \quad R(T_i) \cap W(T_j) \neq \{ \} \quad W(T_i) \cap W(T_j) = \{ \}$$



T_j doesn't read or write anything T_i wrote (but T_i may read something T_j writes).

T_i definitely won't see any of T_j 's writes, because it finishes reading before T_j starts validation, so T_i ordered before T_j .

T_i will always complete its read phase before T_j b/c timestamps assigned after read phase

Applying Conds 1 & 2: Serial Validation

- To validate Xact T:

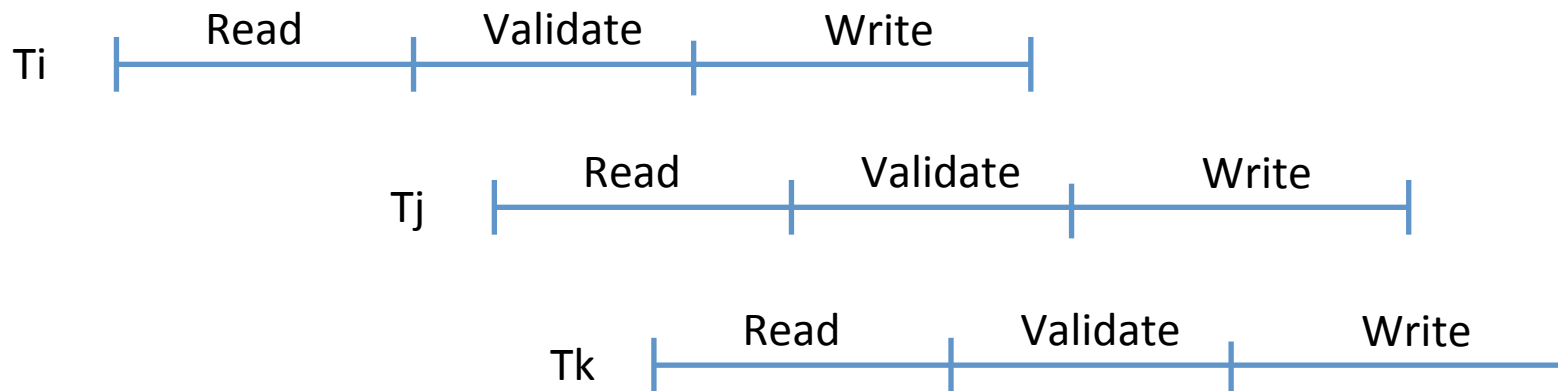
```
valid = true;  
// S = set of Xacts that committed after Begin(T)  
< foreach Ts in S do {  
  if ReadSet(Ts) does not intersect WriteSet(Ts)  
    then valid = false;  
}  
if valid then { install updates; // Write phase  
  Commit T } >  
else Restart T
```

end of critical section



Study Break: Optimistic Concurrency

- What transaction(s) succeed? Which of the three conditions make it possible?



$R(T_i) = d$
 $W(T_i) = a, b, c$

$R(T_j) = d$
 $W(T_j) = b, c$

$R(T_k) = d$
 $W(T_k) = d$

Comments on Serial Validation

- Applies Test 2, with T playing the role of T_j and each $Xact$ in T_s (in turn) being T_i .
- Assignment of $Xact$ id, validation, and the Write phase are inside a **critical section**!
 - I.e., Nothing else goes on concurrently.
 - If Write phase is long, major drawback.
- Optimization for Read-only Xacts:
 - Don't need critical section (because there is no Write phase).

Serial Validation (Contd.)

- **Multistage serial validation:** Validate in stages, at each stage validating T against a subset of the Xacts that committed after $\text{Begin}(T)$.
 - Only last stage has to be inside critical section.
- **Starvation:** Run starving Xact in a critical section (!!)
- **Space for WriteSets:** To validate T_j , must have WriteSets for all T_i where $T_i < T_j$ and T_i was active when T_j began. There may be many such Xacts, and we may run out of space.
 - T_j 's validation fails if it requires a missing WriteSet.

Overhead of Optimistic CC

- Must record read/write activity in ReadSet and WriteSet per Xact.
 - Must create and destroy these sets as needed.
- Must check for conflicts during validation, and must make validated writes ``global'' .
 - Critical section can reduce concurrency.
 - Scheme for making writes global can reduce clustering of objects.
- Optimistic CC restarts Xacts that fail validation.
 - Work done so far is wasted; requires clean-up.

Timestamp Validation

- Maintain timestamps for when each xaction begins and the last time an object was read/written
- Detect conflicts at write time by comparing timestamps
 - Abort as needed

Multiversion Concurrency Control

- To minimize aborts maintain multiple versions of the same object as writes are completed
- Each transaction sees a snapshot of the data at a given time
- If T_i tries to write to something that was read after the transaction started, it will abort

Summary

- Many schemes for managing db concurrency
- Locking is pessimistic – lose time waiting for shared objs
- Optimistic concurrency control detects conflicts when or shortly after they happen
- Best scheme depends on conflict rate
 - OCC best for low conflict – 2PL for mostly overlapping xactions